

In []: project title-~~loan~~ prediction

LINKS RELATED TO PROJECT

f1-scores and accuracies - <https://buffalo.box.com/s/f35i4pfe75kk7y32pesxi3839wnidm69>

VIDEO - <https://buffalo.box.com/s/48apuk09ezfz2ke48hio01gnpk81dqka>

MLFLOW - https://dagshub.com/nithinyanna3/my-first-repo.mlflow/#/experiments/0?searchFilter=&orderByKey=attributes.start_time&orderByAsc=false&startTime=ALL&lifecycleFilter=

DAGSHUB - <https://dagshub.com/nithinyanna3/my-first-repo/experiments>

DOCKER - <https://hub.docker.com/r/nithinyanna/fastapi-app/tags>

Streamlit(cloud deployment) - <https://streamlit-app-p2gr.onrender.com/>

creating a database and loading with the data

```
In [5]: import sqlite3
import pandas as pd
import csv

# Connect to SQLite database (or create it)
conn = sqlite3.connect("loan_db.db")
cursor = conn.cursor()

# Create tables (Person, Loan, Credit, LoanStatus)
cursor.execute('''CREATE TABLE IF NOT EXISTS Person (
    person_id INTEGER PRIMARY KEY,
    person_age INTEGER,
    person_income INTEGER,
    person_home_ownership TEXT,
    person_emp_length INTEGER);''')

cursor.execute('''CREATE TABLE IF NOT EXISTS Loan (
    loan_id INTEGER PRIMARY KEY,
    person_id INTEGER,
    loan_intent TEXT,
    loan_grade TEXT,
    loan_amnt INTEGER,
    loan_int_rate REAL,
    loan_percent_income REAL,
    FOREIGN KEY(person_id) REFERENCES Person(person_id));''')

cursor.execute('''CREATE TABLE IF NOT EXISTS Credit (
    credit_id INTEGER PRIMARY KEY,
    person_id INTEGER,
    cb_person_default_on_file TEXT,
    cb_person_cred_hist_length INTEGER,
    FOREIGN KEY(person_id) REFERENCES Person(person_id));''')
```

```

cursor.execute('''CREATE TABLE IF NOT EXISTS LoanStatus (
                    loan_status_id INTEGER PRIMARY KEY,
                    loan_id INTEGER,
                    loan_status INTEGER,
                    FOREIGN KEY(loan_id) REFERENCES Loan(loan_id));''')

# Parse the CSV file and insert data into tables
with open("C:\\Users\\HP\\Downloads\\train.csv", 'r') as file:
    reader = csv.DictReader(file)
    for row in reader:
        # Insert into Person table
        cursor.execute('''INSERT OR IGNORE INTO Person (person_id, person_age, pers
                        VALUES (?, ?, ?, ?, ?)''',
                        (row['id'], row['person_age'], row['person_income'], row['pe

        # Insert into Loan table
        cursor.execute('''INSERT OR IGNORE INTO Loan (loan_id, person_id, loan_inte
                        VALUES (?, ?, ?, ?, ?, ?, ?)''',
                        (row['id'], row['id'], row['loan_intent'], row['loan_grade']
        cursor.execute('''INSERT OR IGNORE INTO Credit (credit_id, person_id, cb_pe
                        VALUES (?, ?, ?, ?)''',
                        (row['id'], row['id'], row['cb_person_default_on_file'], row

        # Insert into LoanStatus table
        cursor.execute('''INSERT OR IGNORE INTO LoanStatus (loan_status_id, loan_id
                        VALUES (?, ?, ?)''',
                        (row['id'], row['id'], row['loan_status']))

# Commit the changes and close the connection
conn.commit()

# Insert into Credit table

```

joining tables via sql queries

```

In [9]: import sqlite3
import pandas as pd

# Connect to the SQLite database
conn = sqlite3.connect("normalized_db.db")

# SQL query to join the tables and reconstruct the data
query = """
SELECT p.person_id, p.person_age, p.person_income, p.person_home_ownership, p.perso
      l.loan_id, l.loan_intent, l.loan_grade, l.loan_amnt, l.loan_int_rate, l.loan
      c.cb_person_default_on_file, c.cb_person_cred_hist_length, ls.loan_status
FROM Person p
JOIN Loan l ON p.person_id = l.person_id
JOIN Credit c ON p.person_id = c.person_id
JOIN LoanStatus ls ON l.loan_id = ls.loan_id;
"""

# Execute the query and load the result into a Pandas DataFrame
df = pd.read_sql(query, conn)

```

```
# Close the connection  
conn.close()  
  
# Print the resulting DataFrame  
print(df)
```

	person_id	person_age	person_income	person_home_ownership	\
0	0	37	35000	RENT	
1	1	22	56000	OWN	
2	2	29	28800	OWN	
3	3	30	70000	RENT	
4	4	22	60000	RENT	
...	
58640	58640	34	120000	MORTGAGE	
58641	58641	28	28800	RENT	
58642	58642	23	44000	RENT	
58643	58643	22	30000	RENT	
58644	58644	31	75000	MORTGAGE	

	person_emp_length	loan_id	loan_intent	loan_grade	loan_amnt	\
0	0	0	EDUCATION	B	6000	
1	6	1	MEDICAL	C	4000	
2	8	2	PERSONAL	A	6000	
3	14	3	VENTURE	B	12000	
4	2	4	MEDICAL	A	6000	
...	
58640	5	58640	EDUCATION	D	25000	
58641	0	58641	MEDICAL	C	10000	
58642	7	58642	EDUCATION	D	6800	
58643	2	58643	EDUCATION	A	5000	
58644	2	58644	VENTURE	B	15000	

	loan_int_rate	loan_percent_income	cb_person_default_on_file	\
0	11.49	0.17	N	
1	13.35	0.07	N	
2	8.90	0.21	N	
3	11.11	0.17	N	
4	6.92	0.10	N	
...	
58640	15.95	0.21	Y	
58641	12.73	0.35	N	
58642	16.00	0.15	N	
58643	8.90	0.17	N	
58644	11.11	0.20	N	

	cb_person_cred_hist_length	loan_status
0	14	0
1	2	0
2	10	0
3	5	0
4	3	0
...
58640	10	0
58641	8	1
58642	2	1
58643	3	0
58644	5	0

[58645 rows x 14 columns]

removing non data

```
In [12]: df = df.dropna(how='all')
```

removing duplicate rows

```
In [18]: df = df.drop_duplicates()
```

remove rows with missing data NAN

```
In [22]: df = df.dropna(how='any')
```

exploring the dataset

```
In [25]: # Print the entire DataFrame  
print(df)  
  
# Print summary statistics  
print(df.describe())  
  
# Print DataFrame information  
print(df.info())
```

	person_id	person_age	person_income	person_home_ownership	\
0	0	37	35000	RENT	
1	1	22	56000	OWN	
2	2	29	28800	OWN	
3	3	30	70000	RENT	
4	4	22	60000	RENT	
...	
58640	58640	34	120000	MORTGAGE	
58641	58641	28	28800	RENT	
58642	58642	23	44000	RENT	
58643	58643	22	30000	RENT	
58644	58644	31	75000	MORTGAGE	

	person_emp_length	loan_id	loan_intent	loan_grade	loan_amnt	\
0	0	0	EDUCATION	B	6000	
1	6	1	MEDICAL	C	4000	
2	8	2	PERSONAL	A	6000	
3	14	3	VENTURE	B	12000	
4	2	4	MEDICAL	A	6000	
...	
58640	5	58640	EDUCATION	D	25000	
58641	0	58641	MEDICAL	C	10000	
58642	7	58642	EDUCATION	D	6800	
58643	2	58643	EDUCATION	A	5000	
58644	2	58644	VENTURE	B	15000	

	loan_int_rate	loan_percent_income	cb_person_default_on_file	\
0	11.49	0.17	N	
1	13.35	0.07	N	
2	8.90	0.21	N	
3	11.11	0.17	N	
4	6.92	0.10	N	
...	
58640	15.95	0.21	Y	
58641	12.73	0.35	N	
58642	16.00	0.15	N	
58643	8.90	0.17	N	
58644	11.11	0.20	N	

	cb_person_cred_hist_length	loan_status
0	14	0
1	2	0
2	10	0
3	5	0
4	3	0
...
58640	10	0
58641	8	1
58642	2	1
58643	3	0
58644	5	0

[58645 rows x 14 columns]

	person_id	person_age	person_income	person_emp_length	\
count	58645.000000	58645.000000	5.864500e+04	58645.000000	
mean	29322.000000	27.550857	6.404617e+04	4.701015	

std	16929.497605	6.033216	3.793111e+04	3.959784
min	0.000000	20.000000	4.200000e+03	0.000000
25%	14661.000000	23.000000	4.200000e+04	2.000000
50%	29322.000000	26.000000	5.800000e+04	4.000000
75%	43983.000000	30.000000	7.560000e+04	7.000000
max	58644.000000	123.000000	1.900000e+06	123.000000

	loan_id	loan_amnt	loan_int_rate	loan_percent_income \
count	58645.000000	58645.000000	58645.000000	58645.000000
mean	29322.000000	9217.556518	10.677874	0.159238
std	16929.497605	5563.807384	3.034697	0.091692
min	0.000000	500.000000	5.420000	0.000000
25%	14661.000000	5000.000000	7.880000	0.090000
50%	29322.000000	8000.000000	10.750000	0.140000
75%	43983.000000	12000.000000	12.990000	0.210000
max	58644.000000	35000.000000	23.220000	0.830000

	cb_person_cred_hist_length	loan_status
count	58645.000000	58645.000000
mean	5.813556	0.142382
std	4.029196	0.349445
min	2.000000	0.000000
25%	3.000000	0.000000
50%	4.000000	0.000000
75%	8.000000	0.000000
max	30.000000	1.000000

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 58645 entries, 0 to 58644
```

```
Data columns (total 14 columns):
```

#	Column	Non-Null Count	Dtype
0	person_id	58645 non-null	int64
1	person_age	58645 non-null	int64
2	person_income	58645 non-null	int64
3	person_home_ownership	58645 non-null	object
4	person_emp_length	58645 non-null	int64
5	loan_id	58645 non-null	int64
6	loan_intent	58645 non-null	object
7	loan_grade	58645 non-null	object
8	loan_amnt	58645 non-null	int64
9	loan_int_rate	58645 non-null	float64
10	loan_percent_income	58645 non-null	float64
11	cb_person_default_on_file	58645 non-null	object
12	cb_person_cred_hist_length	58645 non-null	int64
13	loan_status	58645 non-null	int64

```
dtypes: float64(2), int64(8), object(4)
```

```
memory usage: 6.3+ MB
```

```
None
```

```
In [29]: print(df.dtypes)
```

```

person_id          int64
person_age         int64
person_income      int64
person_home_ownership  object
person_emp_length  int64
loan_id            int64
loan_intent        object
loan_grade         object
loan_amnt          int64
loan_int_rate      float64
loan_percent_income float64
cb_person_default_on_file  object
cb_person_cred_hist_length int64
loan_status        int64
dtype: object

```

```

In [32]: loan_grade_weights = {'A': 1, 'B': 2, 'C': 3, 'D': 4, 'E': 5}
df['loan_grade'] = df['loan_grade'].map(loan_grade_weights)

# Importance weights for person_home_ownership
home_ownership_weights = {'Own': 1, 'Mortgage': 2, 'Rent': 3}
df['person_home_ownership'] = df['person_home_ownership'].map(home_ownership_weights)

# Importance weights for loan_intent
loan_intent_weights = {'Personal': 3, 'Business': 2, 'Education': 1}
df['loan_intent'] = df['loan_intent'].map(loan_intent_weights)

# Importance weights for cb_person_default_on_file
default_on_file_weights = {'No': 1, 'Yes': 2}
df['cb_person_default_on_file'] = df['cb_person_default_on_file'].map(default_on_file_weights)

# Handling missing or unrecognized categories by filling with default values (0 or 1)
df['loan_grade'] = df['loan_grade'].fillna(0).astype(int)
df['person_home_ownership'] = df['person_home_ownership'].fillna(0).astype(int)
df['loan_intent'] = df['loan_intent'].fillna(0).astype(int)
df['cb_person_default_on_file'] = df['cb_person_default_on_file'].fillna(0).astype(int)

# Display the updated DataFrame
print(df.head())

```


	person_id	person_age	person_income	person_home_ownership	\
0	0	37	35000	0	
1	1	22	56000	0	
2	2	29	28800	0	
3	3	30	70000	0	
4	4	22	60000	0	

	person_emp_length	loan_id	loan_intent	loan_grade	loan_amnt	\
0	0	0	0	2	6000	
1	6	1	0	3	4000	
2	8	2	0	1	6000	
3	14	3	0	2	12000	
4	2	4	0	1	6000	

	loan_int_rate	loan_percent_income	cb_person_default_on_file	\
0	11.49	0.17	0	
1	13.35	0.07	0	
2	8.90	0.21	0	
3	11.11	0.17	0	
4	6.92	0.10	0	

	cb_person_cred_hist_length	loan_status
0	14	0
1	2	0
2	10	0
3	5	0
4	3	0

```
In [35]: import pandas as pd

# Ordinal encoding for loan_grade
loan_grade_weights = {'A': 1, 'B': 2, 'C': 3, 'D': 4, 'E': 5}
df['loan_grade'] = df['loan_grade'].map(loan_grade_weights).fillna(0).astype(int)

# Binary encoding for cb_person_default_on_file
default_on_file_weights = {'N': 0, 'Y': 1}
df['cb_person_default_on_file'] = df['cb_person_default_on_file'].map(default_on_file_weights).fillna(0).astype(int)

# One-Hot Encoding for nominal features
df = pd.get_dummies(df, columns=['person_home_ownership', 'loan_intent'], drop_first=True)

# Display updated DataFrame
print(df.head())
```

	person_id	person_age	person_income	person_emp_length	loan_id	\
0	0	37	35000	0	0	
1	1	22	56000	6	1	
2	2	29	28800	8	2	
3	3	30	70000	14	3	
4	4	22	60000	2	4	

	loan_grade	loan_amnt	loan_int_rate	loan_percent_income	\
0	0	6000	11.49	0.17	
1	0	4000	13.35	0.07	
2	0	6000	8.90	0.21	
3	0	12000	11.11	0.17	
4	0	6000	6.92	0.10	

	cb_person_default_on_file	cb_person_cred_hist_length	loan_status
0	0	14	0
1	0	2	0
2	0	10	0
3	0	5	0
4	0	3	0

saving the data as firstupdate for no confusion

```
In [39]: # Save the updated DataFrame in the Jupyter environment with a new variable name
firstupdate = df.copy()

# Save the new DataFrame to a CSV file
output_path_firstupdate = "C:\\Users\\HP\\Downloads\\firstupdate.csv" # Replace wi
firstupdate.to_csv(output_path_firstupdate, index=False)

print("DataFrame has been saved in the environment as 'firstupdate'")
print(f"DataFrame saved to system as: {output_path_firstupdate}")
print(firstupdate.head()) # Display the first few rows
```

DataFrame has been saved in the environment as 'firstupdate'

DataFrame saved to system as: C:\Users\HP\Downloads\firstupdate.csv

	person_id	person_age	person_income	person_emp_length	loan_id	\
0	0	37	35000	0	0	
1	1	22	56000	6	1	
2	2	29	28800	8	2	
3	3	30	70000	14	3	
4	4	22	60000	2	4	

	loan_grade	loan_amnt	loan_int_rate	loan_percent_income	\
0	0	6000	11.49	0.17	
1	0	4000	13.35	0.07	
2	0	6000	8.90	0.21	
3	0	12000	11.11	0.17	
4	0	6000	6.92	0.10	

	cb_person_default_on_file	cb_person_cred_hist_length	loan_status
0	0	14	0
1	0	2	0
2	0	10	0
3	0	5	0
4	0	3	0

removing the outliers using mahalanobis distance and saving it as second update

```
In [53]: import numpy as np
import pandas as pd
from scipy.spatial import distance
from sklearn.preprocessing import StandardScaler
import matplotlib.pyplot as plt

# Load the DataFrame from environment
df = firstupdate # Assuming 'firstupdate' is already available in the environment

# Select predictors
predictors = df[['person_age', 'person_income', 'person_emp_length', 'loan_amnt',
                 'loan_int_rate', 'loan_percent_income', 'cb_person_cred_hist_length']]

# Normalize the predictors
scaler = StandardScaler()
predictors_scaled = scaler.fit_transform(predictors)

# Calculate the covariance matrix and its inverse
cov_matrix = np.cov(predictors_scaled.T)
inv_cov_matrix = np.linalg.inv(cov_matrix)

# Compute the Mahalanobis distances
mahal_distances = [distance.mahalanobis(x, np.mean(predictors_scaled, axis=0), inv_
                                     for x in predictors_scaled)]

# Add the Mahalanobis distance to the DataFrame
df['mahalanobis'] = mahal_distances

# Identify outliers (using 99th percentile as threshold)
threshold = np.percentile(df['mahalanobis'], 99) # 99th percentile
outliers = df[df['mahalanobis'] > threshold]

# Print the rows that are outliers
print("Outlier rows:")
print(outliers)

# Remove outliers and save the cleaned data to 'secondupdate'
secondupdate = df[df['mahalanobis'] <= threshold].drop(columns=['mahalanobis'])

# Print shapes to confirm
print(f"Original shape: {df.shape}, Shape after removing outliers: {secondupdate.sh

# Optionally, visualize Mahalanobis distances
plt.figure(figsize=(10, 6))
plt.hist(df['mahalanobis'], bins=30, color='skyblue', edgecolor='black')
plt.axvline(threshold, color='red', linestyle='dashed', linewidth=1)
plt.title("Distribution of Mahalanobis Distances")
plt.xlabel("Mahalanobis Distance")
plt.ylabel("Frequency")
plt.show()
```

Outlier rows:

	person_id	person_age	person_income	person_emp_length	loan_id	\
91	91	24	225000	0	91	
583	583	22	250000	4	583	
597	597	60	45000	1	597	
967	967	35	150000	19	967	
1064	1064	46	63000	19	1064	
...	
57844	57844	55	54000	17	57844	
58165	58165	27	156000	9	58165	
58242	58242	53	90000	7	58242	
58436	58436	53	26500	1	58436	
58604	58604	62	150000	3	58604	

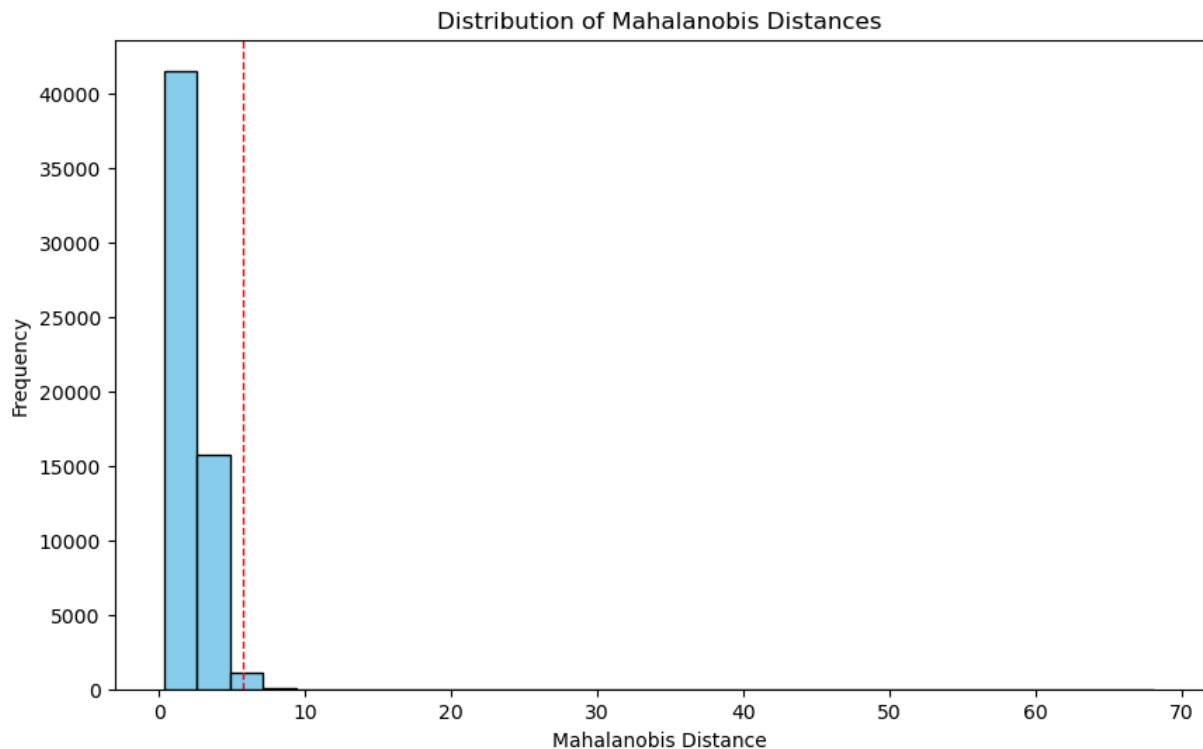
	loan_grade	loan_amnt	loan_int_rate	loan_percent_income	\
91	0	3000	14.26	0.01	
583	0	10000	7.51	0.04	
597	0	8000	12.99	0.18	
967	0	35000	18.39	0.23	
1064	0	7700	18.78	0.12	
...	
57844	0	6000	13.48	0.11	
58165	0	35000	14.27	0.22	
58242	0	5600	8.63	0.06	
58436	0	5000	13.79	0.19	
58604	0	24000	16.95	0.16	

	cb_person_default_on_file	cb_person_cred_hist_length	loan_status	\
91	0	2	0	
583	0	3	0	
597	0	21	0	
967	0	6	0	
1064	0	11	0	
...	
57844	0	18	0	
58165	0	8	0	
58242	0	30	0	
58436	0	30	0	
58604	0	20	0	

	mahalanobis
91	6.154743
583	5.772778
597	6.023917
967	7.165013
1064	6.192444
...	...
57844	5.869117
58165	5.968892
58242	6.482798
58436	6.748854
58604	7.514088

[587 rows x 13 columns]

Original shape: (58645, 13), Shape after removing outliers: (58058, 12)



In [56]: `print(secondupdate.head())`

	person_id	person_age	person_income	person_emp_length	loan_id	\
0	0	37	35000	0	0	
1	1	22	56000	6	1	
2	2	29	28800	8	2	
3	3	30	70000	14	3	
4	4	22	60000	2	4	

	loan_grade	loan_amnt	loan_int_rate	loan_percent_income	\
0	0	6000	11.49	0.17	
1	0	4000	13.35	0.07	
2	0	6000	8.90	0.21	
3	0	12000	11.11	0.17	
4	0	6000	6.92	0.10	

	cb_person_default_on_file	cb_person_cred_hist_length	loan_status
0	0	14	0
1	0	2	0
2	0	10	0
3	0	5	0
4	0	3	0

some exploratory data analysis

```
In [59]: import seaborn as sns
import matplotlib.pyplot as plt

# Univariate Analysis: Histogram, Boxplot, and Density Plot

# List of numeric columns
numeric_columns = ['person_age', 'person_income', 'person_emp_length', 'loan_amnt',
```

```

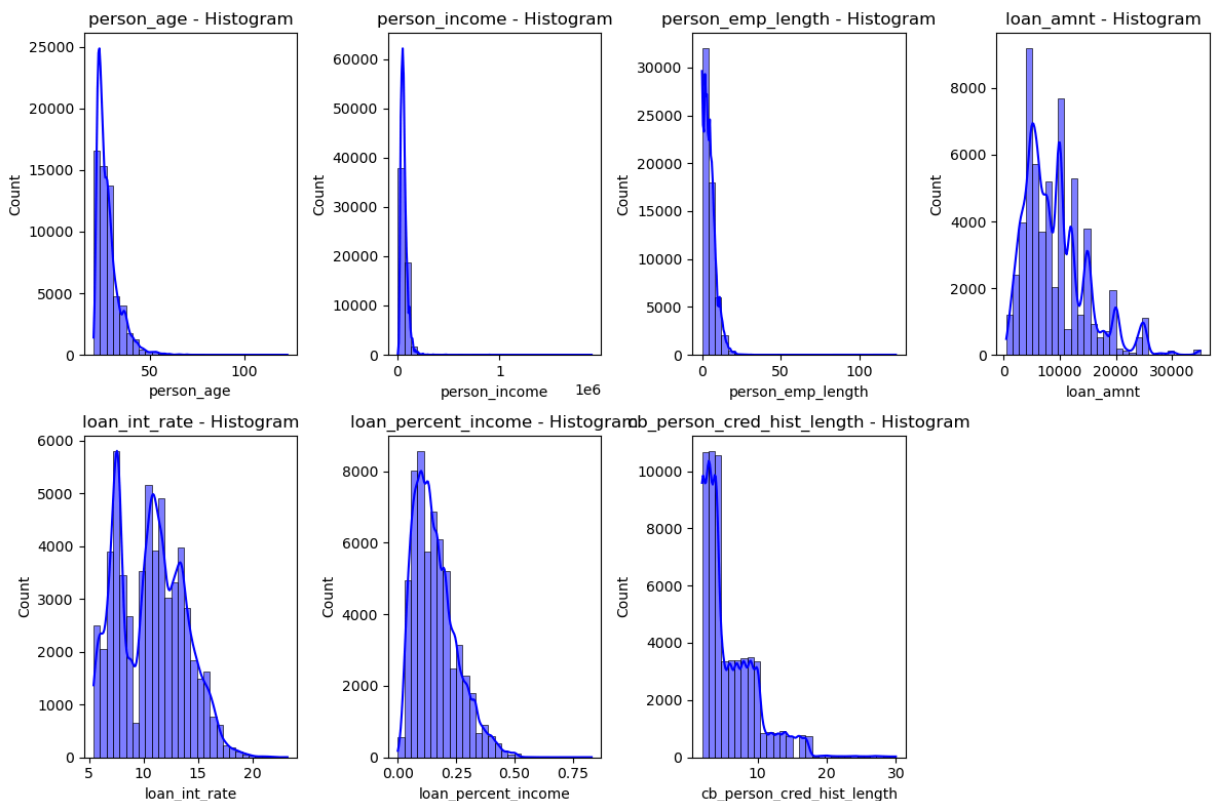
        'loan_int_rate', 'loan_percent_income', 'cb_person_cred_hist_len'

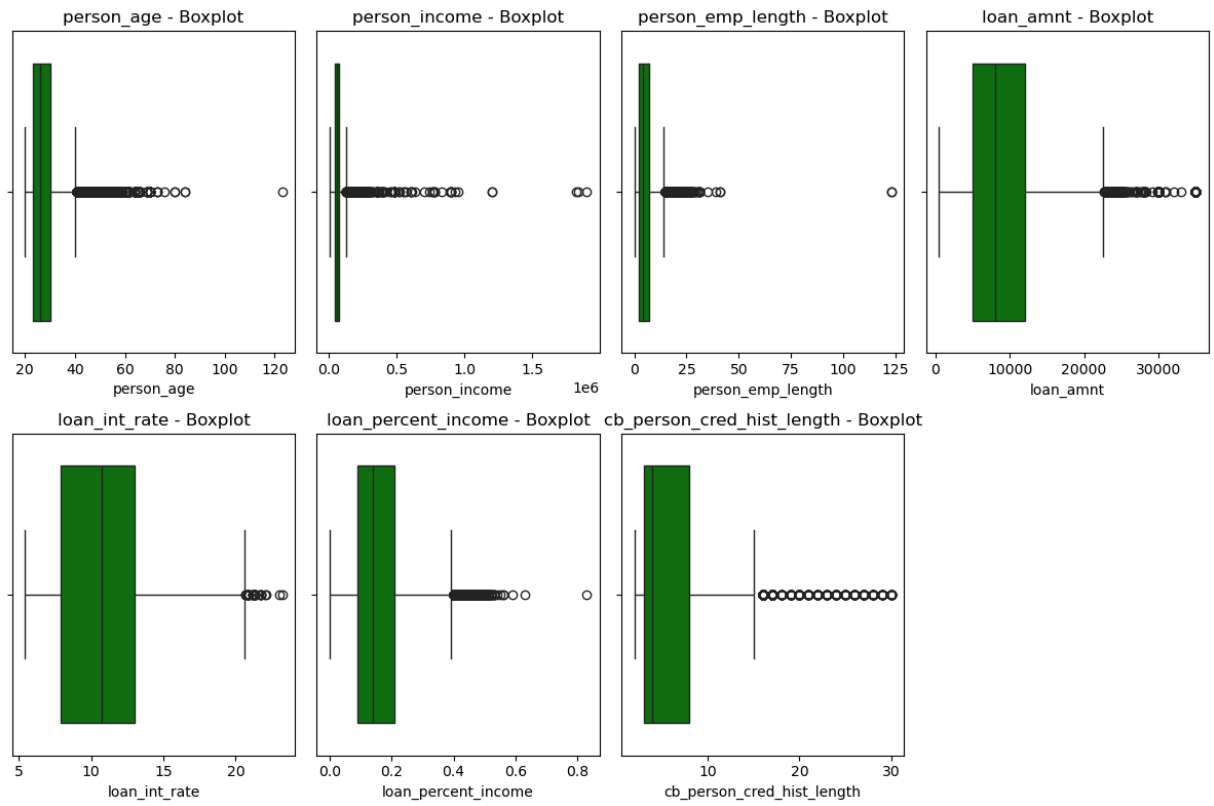
# Plot histograms
plt.figure(figsize=(12, 8))
for i, col in enumerate(numeric_columns, 1):
    plt.subplot(2, 4, i)
    sns.histplot(df[col], kde=True, bins=30, color='blue')
    plt.title(f'{col} - Histogram')
plt.tight_layout()
plt.show()

# Boxplots for Outlier Detection
plt.figure(figsize=(12, 8))
for i, col in enumerate(numeric_columns, 1):
    plt.subplot(2, 4, i)
    sns.boxplot(x=df[col], color='green')
    plt.title(f'{col} - Boxplot')
plt.tight_layout()
plt.show()

# Density Plots
plt.figure(figsize=(12, 8))
for i, col in enumerate(numeric_columns, 1):
    plt.subplot(2, 4, i)
    sns.kdeplot(df[col], shade=True, color='purple')
    plt.title(f'{col} - Density Plot')
plt.tight_layout()
plt.show()

```





```
C:\Users\HP\AppData\Local\Temp\ipykernel_18368\2256259307.py:32: FutureWarning:
`shade` is now deprecated in favor of `fill`; setting `fill=True`.
This will become an error in seaborn v0.14.0; please update your code.

sns.kdeplot(df[col], shade=True, color='purple')
C:\Users\HP\AppData\Local\Temp\ipykernel_18368\2256259307.py:32: FutureWarning:
`shade` is now deprecated in favor of `fill`; setting `fill=True`.
This will become an error in seaborn v0.14.0; please update your code.

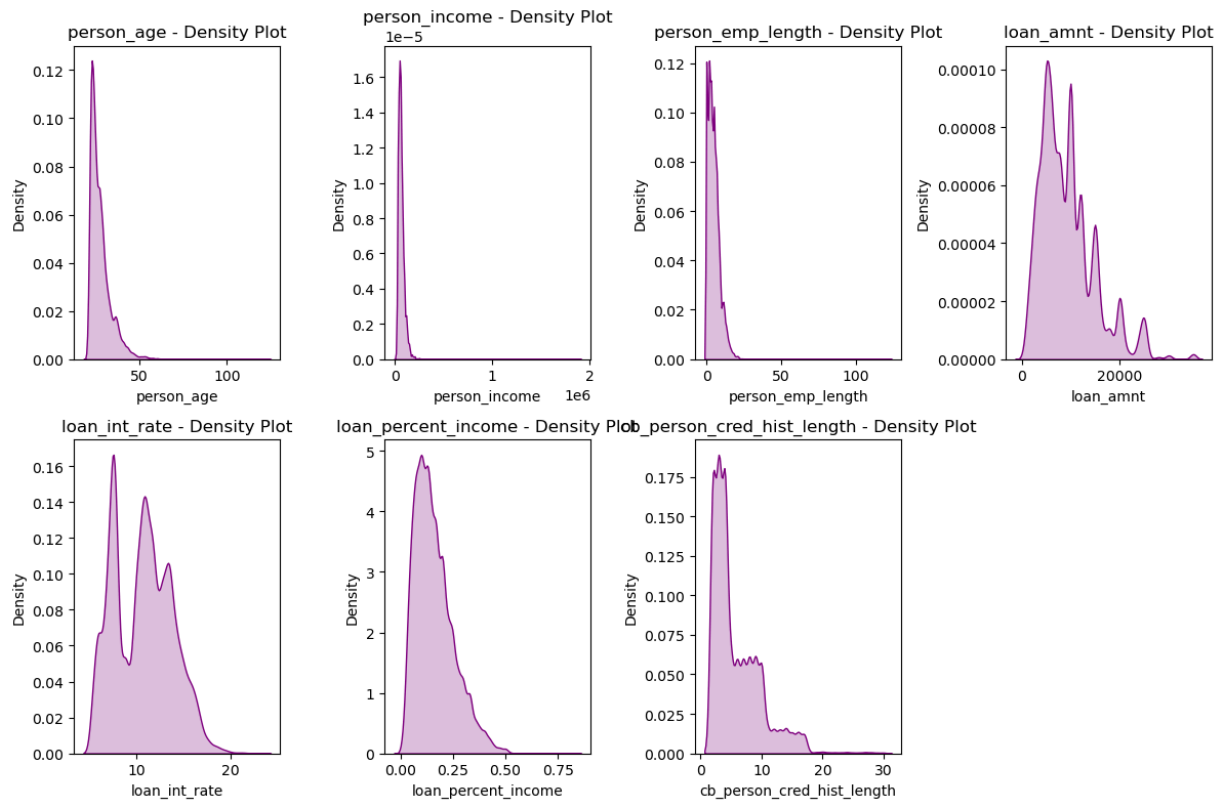
sns.kdeplot(df[col], shade=True, color='purple')
C:\Users\HP\AppData\Local\Temp\ipykernel_18368\2256259307.py:32: FutureWarning:
`shade` is now deprecated in favor of `fill`; setting `fill=True`.
This will become an error in seaborn v0.14.0; please update your code.

sns.kdeplot(df[col], shade=True, color='purple')
C:\Users\HP\AppData\Local\Temp\ipykernel_18368\2256259307.py:32: FutureWarning:
`shade` is now deprecated in favor of `fill`; setting `fill=True`.
This will become an error in seaborn v0.14.0; please update your code.

sns.kdeplot(df[col], shade=True, color='purple')
C:\Users\HP\AppData\Local\Temp\ipykernel_18368\2256259307.py:32: FutureWarning:
`shade` is now deprecated in favor of `fill`; setting `fill=True`.
This will become an error in seaborn v0.14.0; please update your code.

sns.kdeplot(df[col], shade=True, color='purple')
C:\Users\HP\AppData\Local\Temp\ipykernel_18368\2256259307.py:32: FutureWarning:
`shade` is now deprecated in favor of `fill`; setting `fill=True`.
This will become an error in seaborn v0.14.0; please update your code.

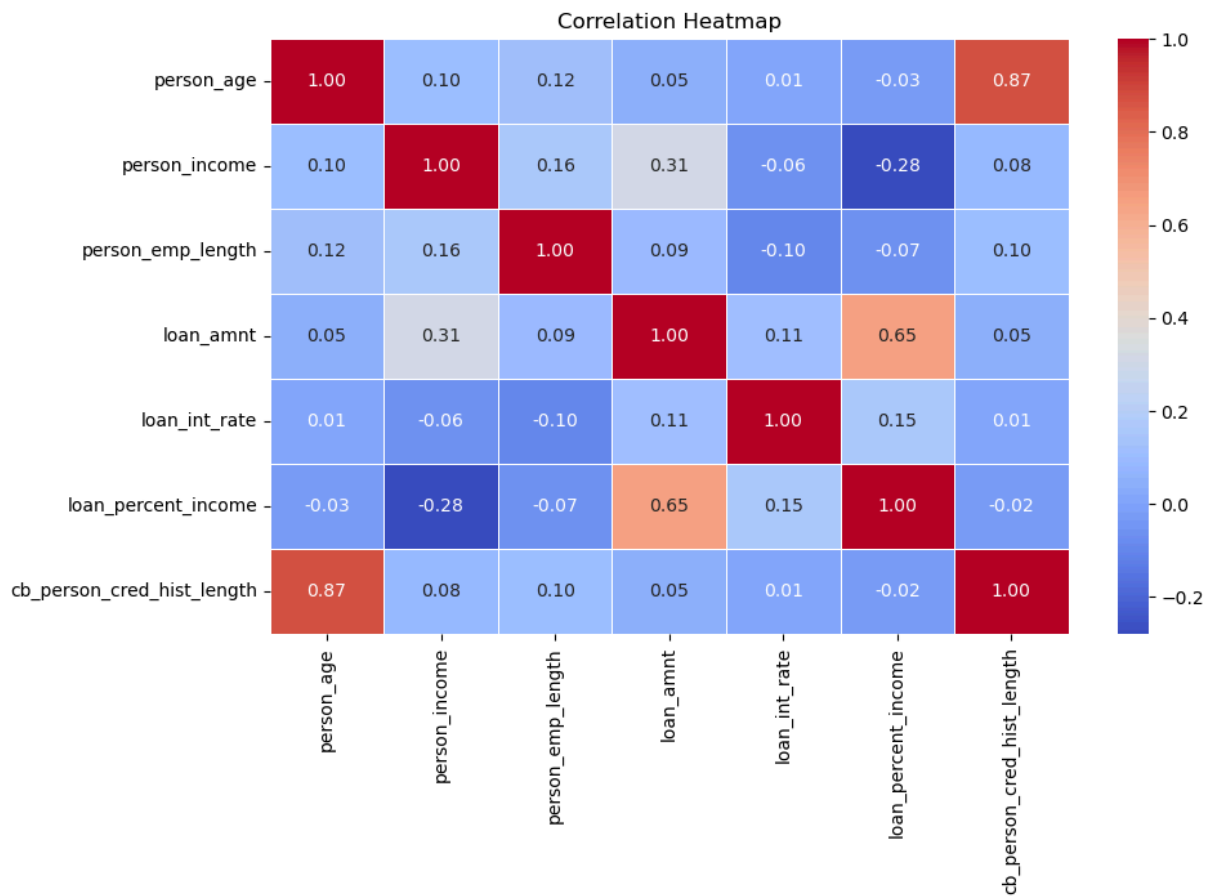
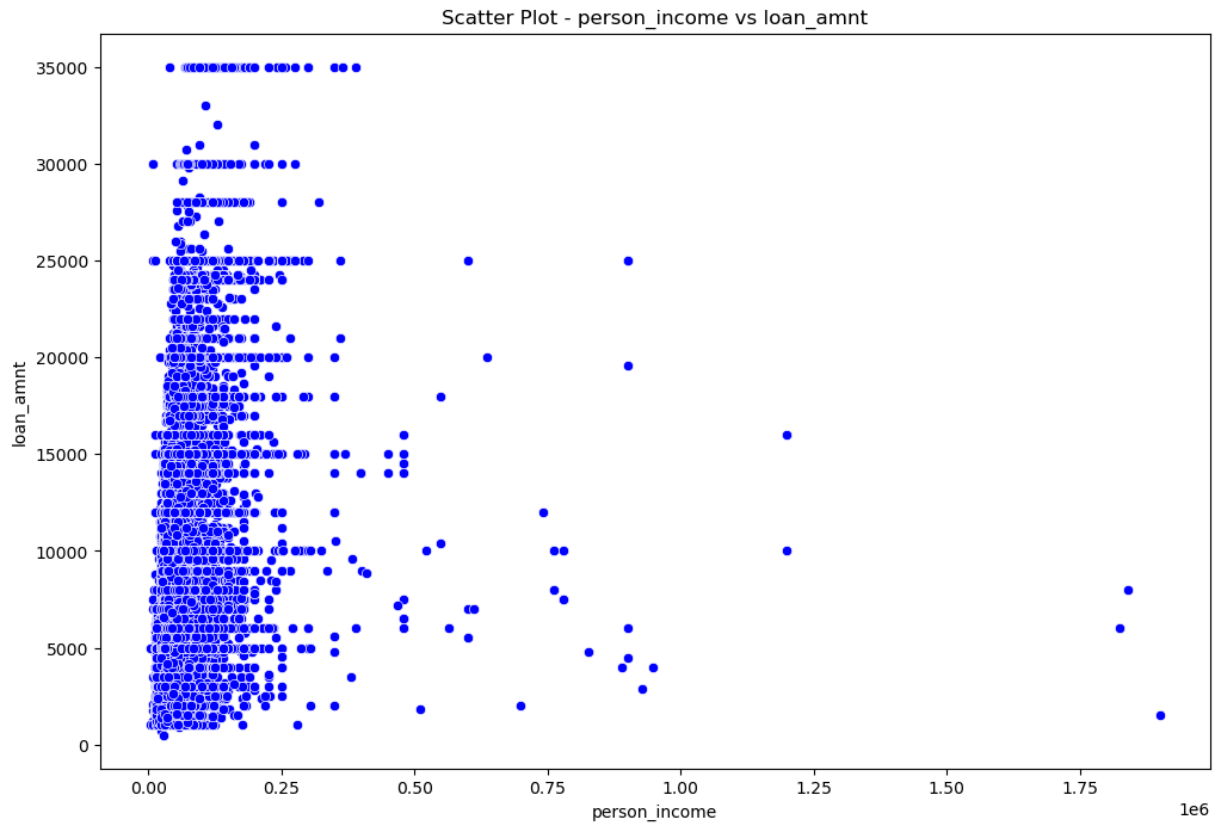
sns.kdeplot(df[col], shade=True, color='purple')
```

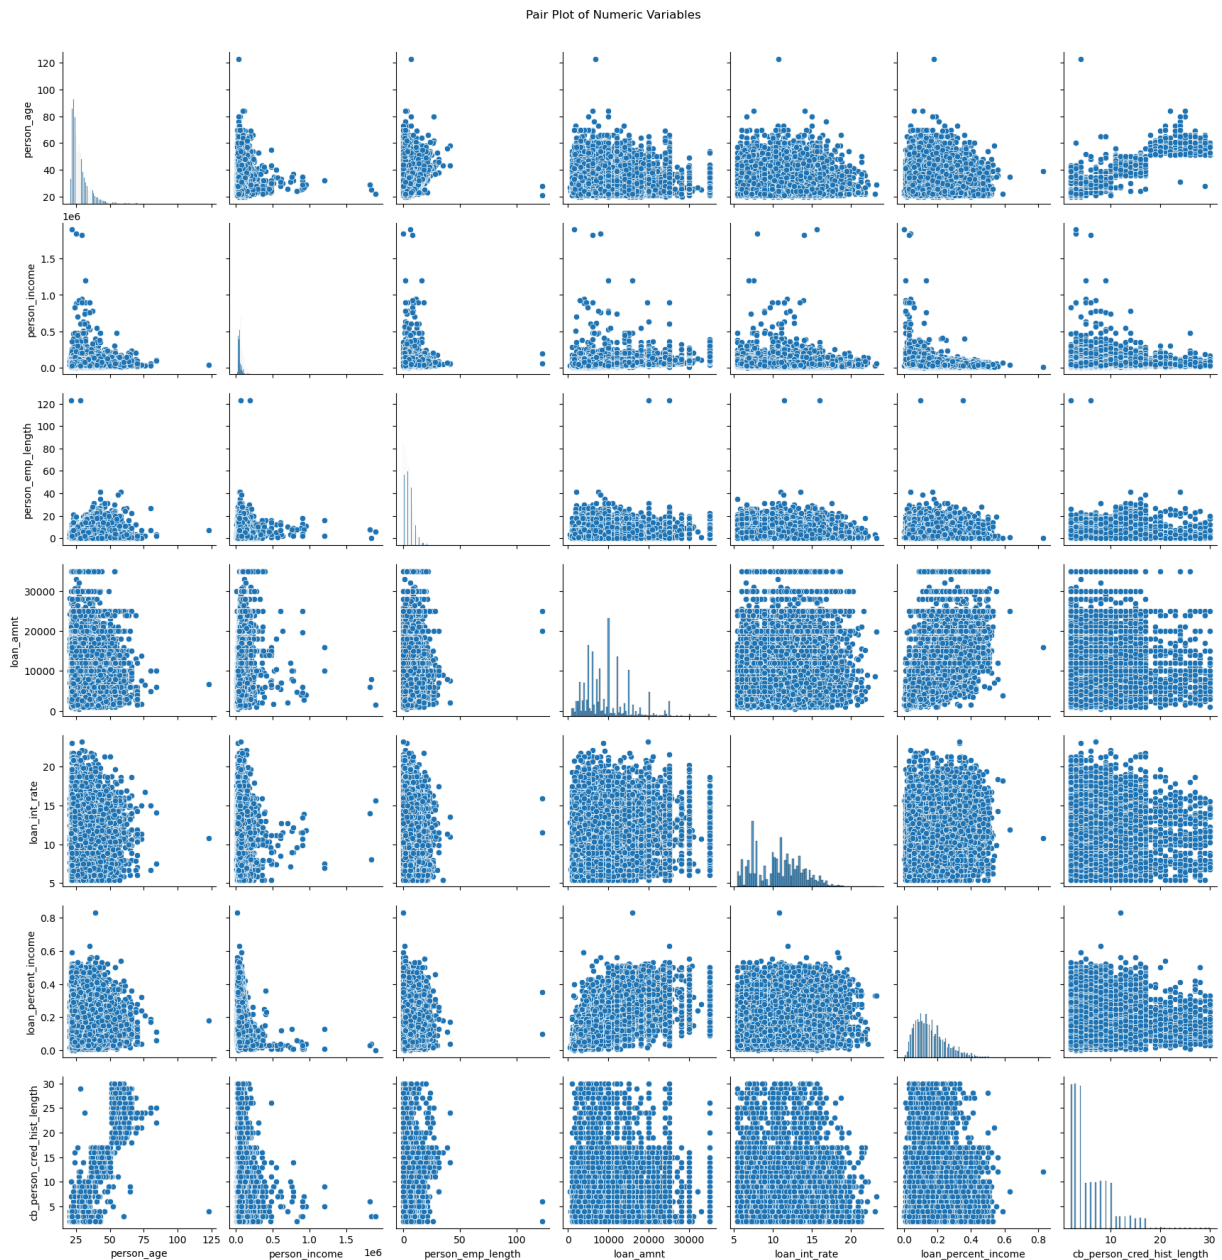



```
In [62]: # Scatter Plots
plt.figure(figsize=(12, 8))
sns.scatterplot(x='person_income', y='loan_amnt', data=df, color='blue')
plt.title('Scatter Plot - person_income vs loan_amnt')
plt.show()

# Correlation Heatmap
corr_matrix = df[numeric_columns].corr()
plt.figure(figsize=(10, 6))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', fmt='.2f', linewidths=0.5)
plt.title('Correlation Heatmap')
plt.show()

# Pair Plot
sns.pairplot(df[numeric_columns], height=2.5)
plt.suptitle('Pair Plot of Numeric Variables', y=1.02)
plt.show()
```





```
In [66]: from mpl_toolkits.mplot3d import Axes3D
from sklearn.preprocessing import LabelEncoder
import pandas as pd

# Multivariate Analysis: 3D Plot

fig = plt.figure(figsize=(10, 8))
ax = fig.add_subplot(111, projection='3d')
ax.scatter(df['person_age'], df['loan_amnt'], df['loan_int_rate'], c='purple', marker='o')
ax.set_xlabel('person_age')
ax.set_ylabel('loan_amnt')
ax.set_zlabel('loan_int_rate')
plt.title('3D Scatter Plot - person_age, loan_amnt, and loan_int_rate')
plt.show()

# Parallel Coordinate Plot (requires normalization)
from pandas.plotting import parallel_coordinates
```

```


# Normalizing data for parallel coordinates plot
df_normalized = df[numeric_columns].copy()
df_normalized = (df_normalized - df_normalized.min()) / (df_normalized.max() - df_normalized.min())

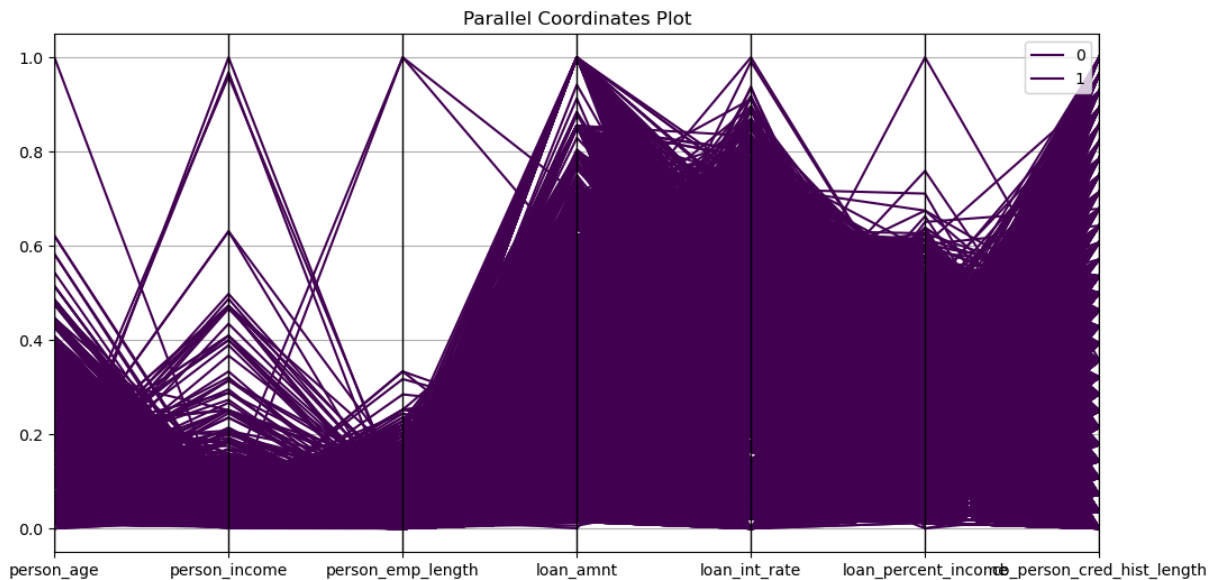
# Adding a target column for the color differentiation (using loan_status as an example)
df_normalized['loan_status'] = df['loan_status']

# Plotting the parallel coordinate plot
plt.figure(figsize=(12, 6))
parallel_coordinates(df_normalized, 'loan_status', color=plt.cm.viridis(df_normalized['loan_status']))
plt.title('Parallel Coordinates Plot')
plt.show()

```

3D Scatter Plot - person_age, loan_amnt, and loan_int_rate





transforming the test data

loading the test data and exploring

In [76]: `import pandas as pd`

Load the test data

`test_data = pd.read_csv('C:\\Users\\HP\\Downloads\\test.csv')`

Display the head of the test data

`print(test_data.head())`

	id	person_age	person_income	person_home_ownership	person_emp_length	\
0	58645	23	69000	RENT	3.0	
1	58646	26	96000	MORTGAGE	6.0	
2	58647	26	30000	RENT	5.0	
3	58648	33	50000	RENT	4.0	
4	58649	26	102000	MORTGAGE	8.0	

	loan_intent	loan_grade	loan_amnt	loan_int_rate	\
0	HOMEIMPROVEMENT	F	25000	15.76	
1	PERSONAL	C	10000	12.68	
2	VENTURE	E	4000	17.19	
3	DEBTCONSOLIDATION	A	7000	8.90	
4	HOMEIMPROVEMENT	D	15000	16.32	

	loan_percent_income	cb_person_default_on_file	cb_person_cred_hist_length
0	0.36	N	2
1	0.10	Y	4
2	0.13	Y	2
3	0.14	N	7
4	0.15	Y	4

converting the test data into the form of training data for predictions

In [86]: `import pandas as pd`

```

# Load the test data
test_data = pd.read_csv('C:\\Users\\HP\\Downloads\\test.csv')

# Drop the 'person_home_ownership' column
test_data = test_data.drop(columns=['person_home_ownership', 'loan_intent', 'id'],)

# Label encoding for required columns
home_ownership_weights = {'Own': 1, 'Mortgage': 2, 'Rent': 3}
loan_grade_weights = {'A': 0, 'B': 1, 'C': 2, 'D': 3, 'E': 4, 'F': 5}

test_data['loan_grade'] = test_data['loan_grade'].map(loan_grade_weights)

# Convert 'cb_person_default_on_file' from strings to numeric
default_on_file_weights = {'N': 1, 'Y': 2}
test_data['cb_person_default_on_file'] = test_data['cb_person_default_on_file'].map

# Save the updated test data to the environment
test_data.to_pickle('testdata.pkl')

# Optionally, save it as a CSV
test_data.to_csv('updatedtestdata.csv', index=False)

# Display the head of the updated test data
print(test_data.head())

```

	person_age	person_income	person_emp_length	loan_grade	loan_amnt	\
0	23	69000	3.0	5.0	25000	
1	26	96000	6.0	2.0	10000	
2	26	30000	5.0	4.0	4000	
3	33	50000	4.0	0.0	7000	
4	26	102000	8.0	3.0	15000	

	loan_int_rate	loan_percent_income	cb_person_default_on_file	\
0	15.76	0.36	1	
1	12.68	0.10	2	
2	17.19	0.13	2	
3	8.90	0.14	1	
4	16.32	0.15	2	

	cb_person_cred_hist_length
0	2
1	4
2	2
3	7
4	4

```

In [94]: import pandas as pd

columns_to_drop = ['person_id', 'loan_id'] # Add more columns if necessary
secondupdate = secondupdate.drop(columns=columns_to_drop)

# Print the head of the updated DataFrame
print(secondupdate.head())

```

	person_age	person_income	person_emp_length	loan_grade	loan_amnt	\
0	37	35000	0	0	6000	
1	22	56000	6	0	4000	
2	29	28800	8	0	6000	
3	30	70000	14	0	12000	
4	22	60000	2	0	6000	

	loan_int_rate	loan_percent_income	cb_person_default_on_file	\
0	11.49	0.17	0	
1	13.35	0.07	0	
2	8.90	0.21	0	
3	11.11	0.17	0	
4	6.92	0.10	0	

	cb_person_cred_hist_length	loan_status
0	14	0
1	2	0
2	10	0
3	5	0
4	3	0

importing dagshub for logging ml model details

In [100...

```
pip install mlflow dagshub
```

Requirement already satisfied: mlflow in c:\users\hp\anaconda3\lib\site-packages (2.19.0)

Requirement already satisfied: dagshub in c:\users\hp\anaconda3\lib\site-packages (0.3.47)

Requirement already satisfied: mlflow-skinny==2.19.0 in c:\users\hp\anaconda3\lib\site-packages (from mlflow) (2.19.0)

Requirement already satisfied: Flask<4 in c:\users\hp\anaconda3\lib\site-packages (from mlflow) (3.0.3)

Requirement already satisfied: Jinja2<4,>=3.0 in c:\users\hp\anaconda3\lib\site-packages (from mlflow) (3.1.4)

Requirement already satisfied: alembic!=1.10.0,<2 in c:\users\hp\anaconda3\lib\site-packages (from mlflow) (1.14.0)

Requirement already satisfied: docker<8,>=4.0.0 in c:\users\hp\anaconda3\lib\site-packages (from mlflow) (7.1.0)

Requirement already satisfied: graphene<4 in c:\users\hp\anaconda3\lib\site-packages (from mlflow) (3.4.3)

Requirement already satisfied: markdown<4,>=3.3 in c:\users\hp\anaconda3\lib\site-packages (from mlflow) (3.4.1)

Requirement already satisfied: matplotlib<4 in c:\users\hp\anaconda3\lib\site-packages (from mlflow) (3.8.4)

Requirement already satisfied: numpy<3 in c:\users\hp\anaconda3\lib\site-packages (from mlflow) (1.26.4)

Requirement already satisfied: pandas<3 in c:\users\hp\anaconda3\lib\site-packages (from mlflow) (2.2.2)

Requirement already satisfied: pyarrow<19,>=4.0.0 in c:\users\hp\anaconda3\lib\site-packages (from mlflow) (14.0.2)

Requirement already satisfied: scikit-learn<2 in c:\users\hp\anaconda3\lib\site-packages (from mlflow) (1.4.2)

Requirement already satisfied: scipy<2 in c:\users\hp\anaconda3\lib\site-packages (from mlflow) (1.13.1)

Requirement already satisfied: sqlalchemy<3,>=1.4.0 in c:\users\hp\anaconda3\lib\site-packages (from mlflow) (2.0.30)

Requirement already satisfied: waitress<4 in c:\users\hp\anaconda3\lib\site-packages (from mlflow) (3.0.2)

Requirement already satisfied: cachetools<6,>=5.0.0 in c:\users\hp\anaconda3\lib\site-packages (from mlflow-skinny==2.19.0->mlflow) (5.3.3)

Requirement already satisfied: click<9,>=7.0 in c:\users\hp\anaconda3\lib\site-packages (from mlflow-skinny==2.19.0->mlflow) (8.1.7)

Requirement already satisfied: cloudpickle<4 in c:\users\hp\anaconda3\lib\site-packages (from mlflow-skinny==2.19.0->mlflow) (2.2.1)

Requirement already satisfied: databricks-sdk<1,>=0.20.0 in c:\users\hp\anaconda3\lib\site-packages (from mlflow-skinny==2.19.0->mlflow) (0.39.0)

Requirement already satisfied: gitpython<4,>=3.1.9 in c:\users\hp\anaconda3\lib\site-packages (from mlflow-skinny==2.19.0->mlflow) (3.1.37)

Requirement already satisfied: importlib_metadata!=4.7.0,<9,>=3.7.0 in c:\users\hp\anaconda3\lib\site-packages (from mlflow-skinny==2.19.0->mlflow) (7.0.1)

Requirement already satisfied: opentelemetry-api<3,>=1.9.0 in c:\users\hp\anaconda3\lib\site-packages (from mlflow-skinny==2.19.0->mlflow) (1.29.0)

Requirement already satisfied: opentelemetry-sdk<3,>=1.9.0 in c:\users\hp\anaconda3\lib\site-packages (from mlflow-skinny==2.19.0->mlflow) (1.29.0)

Requirement already satisfied: packaging<25 in c:\users\hp\anaconda3\lib\site-packages (from mlflow-skinny==2.19.0->mlflow) (23.2)

Requirement already satisfied: protobuf<6,>=3.12.0 in c:\users\hp\anaconda3\lib\site-packages (from mlflow-skinny==2.19.0->mlflow) (3.20.3)

Requirement already satisfied: pyyaml<7,>=5.1 in c:\users\hp\anaconda3\lib\site-packages (from mlflow-skinny==2.19.0->mlflow) (6.0.1)

Requirement already satisfied: requests<3,>=2.17.3 in c:\users\hp\anaconda3\lib\site-packages (from mlflow-skinny==2.19.0->mlflow) (2.32.2)

Requirement already satisfied: sqlparse<1,>=0.4.0 in c:\users\hp\anaconda3\lib\site-packages (from mlflow-skinny==2.19.0->mlflow) (0.5.3)

Requirement already satisfied: appdirs>=1.4.4 in c:\users\hp\anaconda3\lib\site-packages (from dagshub) (1.4.4)

Requirement already satisfied: httpx>=0.23.0 in c:\users\hp\anaconda3\lib\site-packages (from dagshub) (0.27.0)

Requirement already satisfied: rich>=13.1.0 in c:\users\hp\anaconda3\lib\site-packages (from dagshub) (13.3.5)

Requirement already satisfied: dacite~=1.6.0 in c:\users\hp\anaconda3\lib\site-packages (from dagshub) (1.6.0)

Requirement already satisfied: tenacity>=8.2.2 in c:\users\hp\anaconda3\lib\site-packages (from dagshub) (8.2.2)

Requirement already satisfied: gql[requests] in c:\users\hp\anaconda3\lib\site-packages (from dagshub) (3.5.0)

Requirement already satisfied: dataclasses-json in c:\users\hp\anaconda3\lib\site-packages (from dagshub) (0.6.7)

Requirement already satisfied: treelib>=1.6.4 in c:\users\hp\anaconda3\lib\site-packages (from dagshub) (1.7.0)

Requirement already satisfied: pathvalidate>=3.0.0 in c:\users\hp\anaconda3\lib\site-packages (from dagshub) (3.2.1)

Requirement already satisfied: python-dateutil in c:\users\hp\anaconda3\lib\site-packages (from dagshub) (2.9.0.post0)

Requirement already satisfied: boto3 in c:\users\hp\anaconda3\lib\site-packages (from dagshub) (1.35.81)

Requirement already satisfied: dagshub-annotation-converter>=0.1.0 in c:\users\hp\anaconda3\lib\site-packages (from dagshub) (0.1.2)

Requirement already satisfied: Mako in c:\users\hp\anaconda3\lib\site-packages (from alembic!=1.10.0,<2->mlflow) (1.3.8)

Requirement already satisfied: typing-extensions>=4 in c:\users\hp\anaconda3\lib\site-packages (from alembic!=1.10.0,<2->mlflow) (4.11.0)

Requirement already satisfied: colorama in c:\users\hp\anaconda3\lib\site-packages (from click<9,>=7.0->mlflow-skinny==2.19.0->mlflow) (0.4.6)

Requirement already satisfied: lxml in c:\users\hp\anaconda3\lib\site-packages (from dagshub-annotation-converter>=0.1.0->dagshub) (5.2.1)

Requirement already satisfied: pillow in c:\users\hp\anaconda3\lib\site-packages (from dagshub-annotation-converter>=0.1.0->dagshub) (10.3.0)

Requirement already satisfied: pydantic>=2.0.0 in c:\users\hp\anaconda3\lib\site-packages (from dagshub-annotation-converter>=0.1.0->dagshub) (2.5.3)

Requirement already satisfied: pywin32>=304 in c:\users\hp\anaconda3\lib\site-packages (from docker<8,>=4.0.0->mlflow) (305.1)

Requirement already satisfied: urllib3>=1.26.0 in c:\users\hp\anaconda3\lib\site-packages (from docker<8,>=4.0.0->mlflow) (2.2.2)

Requirement already satisfied: Werkzeug>=3.0.0 in c:\users\hp\anaconda3\lib\site-packages (from Flask<4->mlflow) (3.0.3)

Requirement already satisfied: itsdangerous>=2.1.2 in c:\users\hp\anaconda3\lib\site-packages (from Flask<4->mlflow) (2.2.0)

Requirement already satisfied: blinker>=1.6.2 in c:\users\hp\anaconda3\lib\site-packages (from Flask<4->mlflow) (1.6.2)

Requirement already satisfied: gitdb<5,>=4.0.1 in c:\users\hp\anaconda3\lib\site-packages (from gitpython<4,>=3.1.9->mlflow-skinny==2.19.0->mlflow) (4.0.7)

Requirement already satisfied: graphql-core<3.3,>=3.1 in c:\users\hp\anaconda3\lib\site-packages (from graphene<4->mlflow) (3.2.5)

Requirement already satisfied: graphql-relay<3.3,>=3.1 in c:\users\hp\anaconda3\lib\site-packages (from graphene<4->mlflow) (3.2.0)

Requirement already satisfied: anyio in c:\users\hp\anaconda3\lib\site-packages (from httpx>=0.23.0->dagshub) (4.2.0)

Requirement already satisfied: certifi in c:\users\hp\anaconda3\lib\site-packages (from httpx>=0.23.0->dagshub) (2024.7.4)

Requirement already satisfied: httpcore==1.* in c:\users\hp\anaconda3\lib\site-packages (from httpx>=0.23.0->dagshub) (1.0.2)

Requirement already satisfied: idna in c:\users\hp\anaconda3\lib\site-packages (from httpx>=0.23.0->dagshub) (3.7)

Requirement already satisfied: sniffio in c:\users\hp\anaconda3\lib\site-packages (from httpx>=0.23.0->dagshub) (1.3.0)

Requirement already satisfied: h11<0.15,>=0.13 in c:\users\hp\anaconda3\lib\site-packages (from httpcore==1.*->httpx>=0.23.0->dagshub) (0.14.0)

Requirement already satisfied: MarkupSafe>=2.0 in c:\users\hp\anaconda3\lib\site-packages (from Jinja2<4,>=3.0->mlflow) (2.1.3)

Requirement already satisfied: contourpy>=1.0.1 in c:\users\hp\anaconda3\lib\site-packages (from matplotlib<4->mlflow) (1.2.0)

Requirement already satisfied: cycler>=0.10 in c:\users\hp\anaconda3\lib\site-packages (from matplotlib<4->mlflow) (0.11.0)

Requirement already satisfied: fonttools>=4.22.0 in c:\users\hp\anaconda3\lib\site-packages (from matplotlib<4->mlflow) (4.51.0)

Requirement already satisfied: kiwisolver>=1.3.1 in c:\users\hp\anaconda3\lib\site-packages (from matplotlib<4->mlflow) (1.4.4)

Requirement already satisfied: pyparsing>=2.3.1 in c:\users\hp\anaconda3\lib\site-packages (from matplotlib<4->mlflow) (3.0.9)

Requirement already satisfied: pytz>=2020.1 in c:\users\hp\anaconda3\lib\site-packages (from pandas<3->mlflow) (2024.1)

Requirement already satisfied: tzdata>=2022.7 in c:\users\hp\anaconda3\lib\site-packages (from pandas<3->mlflow) (2023.3)

Requirement already satisfied: six>=1.5 in c:\users\hp\anaconda3\lib\site-packages (from python-dateutil->dagshub) (1.16.0)

Requirement already satisfied: markdown-it-py<3.0.0,>=2.2.0 in c:\users\hp\anaconda3\lib\site-packages (from rich>=13.1.0->dagshub) (2.2.0)

Requirement already satisfied: pygments<3.0.0,>=2.13.0 in c:\users\hp\anaconda3\lib\site-packages (from rich>=13.1.0->dagshub) (2.15.1)

Requirement already satisfied: joblib>=1.2.0 in c:\users\hp\anaconda3\lib\site-packages (from scikit-learn<2->mlflow) (1.4.2)

Requirement already satisfied: threadpoolctl>=2.0.0 in c:\users\hp\anaconda3\lib\site-packages (from scikit-learn<2->mlflow) (2.2.0)

Requirement already satisfied: greenlet!=0.4.17 in c:\users\hp\anaconda3\lib\site-packages (from sqlalchemy<3,>=1.4.0->mlflow) (3.0.1)

Requirement already satisfied: botocore<1.36.0,>=1.35.81 in c:\users\hp\anaconda3\lib\site-packages (from boto3->dagshub) (1.35.81)

Requirement already satisfied: jmespath<2.0.0,>=0.7.1 in c:\users\hp\anaconda3\lib\site-packages (from boto3->dagshub) (1.0.1)

Requirement already satisfied: s3transfer<0.11.0,>=0.10.0 in c:\users\hp\anaconda3\lib\site-packages (from boto3->dagshub) (0.10.4)

Requirement already satisfied: marshmallow<4.0.0,>=3.18.0 in c:\users\hp\anaconda3\lib\site-packages (from dataclasses-json->dagshub) (3.23.1)

Requirement already satisfied: typing-inspect<1,>=0.4.0 in c:\users\hp\anaconda3\lib\site-packages (from dataclasses-json->dagshub) (0.9.0)

Requirement already satisfied: yarl<2.0,>=1.6 in c:\users\hp\anaconda3\lib\site-packages (from gql[requests]->dagshub) (1.9.3)

Requirement already satisfied: backoff<3.0,>=1.11.1 in c:\users\hp\anaconda3\lib\site-packages (from gql[requests]->dagshub) (2.2.1)

Requirement already satisfied: requests-toolbelt<2,>=1.0.0 in c:\users\hp\anaconda3\lib\site-packages (from gql[requests]->dagshub) (1.0.0)

Requirement already satisfied: google-auth~=2.0 in c:\users\hp\anaconda3\lib\site-packages (from databricks-sdk<1,>=0.20.0->mlflow-skinny==2.19.0->mlflow) (2.37.0)

Requirement already satisfied: smmap<5,>=3.0.1 in c:\users\hp\anaconda3\lib\site-packages (from gitdb<5,>=4.0.1->gitpython<4,>=3.1.9->mlflow-skinny==2.19.0->mlflow) (4.0.0)

Requirement already satisfied: zipp>=0.5 in c:\users\hp\anaconda3\lib\site-packages (from importlib_metadata!=4.7.0,<9,>=3.7.0->mlflow-skinny==2.19.0->mlflow) (3.17.0)

Requirement already satisfied: mdurl~=0.1 in c:\users\hp\anaconda3\lib\site-packages (from markdown-it-py<3.0.0,>=2.2.0->rich>=13.1.0->dagshub) (0.1.0)

Requirement already satisfied: deprecated>=1.2.6 in c:\users\hp\anaconda3\lib\site-packages (from opentelemetry-api<3,>=1.9.0->mlflow-skinny==2.19.0->mlflow) (1.2.15)

Requirement already satisfied: opentelemetry-semantic-conventions==0.50b0 in c:\users\hp\anaconda3\lib\site-packages (from opentelemetry-sdk<3,>=1.9.0->mlflow-skinny==2.19.0->mlflow) (0.50b0)

Requirement already satisfied: annotated-types>=0.4.0 in c:\users\hp\anaconda3\lib\site-packages (from pydantic>=2.0.0->dagshub-annotation-converter>=0.1.0->dagshub) (0.6.0)

Requirement already satisfied: pydantic-core==2.14.6 in c:\users\hp\anaconda3\lib\site-packages (from pydantic>=2.0.0->dagshub-annotation-converter>=0.1.0->dagshub) (2.14.6)

Requirement already satisfied: charset-normalizer<4,>=2 in c:\users\hp\anaconda3\lib\site-packages (from requests<3,>=2.17.3->mlflow-skinny==2.19.0->mlflow) (2.0.4)

Requirement already satisfied: mpy-extensions>=0.3.0 in c:\users\hp\anaconda3\lib\site-packages (from typing-inspect<1,>=0.4.0->dataclasses-json->dagshub) (1.0.0)

Requirement already satisfied: multidict>=4.0 in c:\users\hp\anaconda3\lib\site-packages (from yarl<2.0,>=1.6->gql[requests]->dagshub) (6.0.4)

Requirement already satisfied: wrapt<2,>=1.10 in c:\users\hp\anaconda3\lib\site-packages (from deprecated>=1.2.6->opentelemetry-api<3,>=1.9.0->mlflow-skinny==2.19.0->mlflow) (1.14.1)

Requirement already satisfied: pyasn1-modules>=0.2.1 in c:\users\hp\anaconda3\lib\site-packages (from google-auth~=2.0->databricks-sdk<1,>=0.20.0->mlflow-skinny==2.19.0->mlflow) (0.2.8)

Requirement already satisfied: rsa<5,>=3.1.4 in c:\users\hp\anaconda3\lib\site-packages (from google-auth~=2.0->databricks-sdk<1,>=0.20.0->mlflow-skinny==2.19.0->mlflow) (4.9)

Requirement already satisfied: pyasn1<0.5.0,>=0.4.6 in c:\users\hp\anaconda3\lib\site-packages (from pyasn1-modules>=0.2.1->google-auth~=2.0->databricks-sdk<1,>=0.20.0->mlflow-skinny==2.19.0->mlflow) (0.4.8)

Note: you may need to restart the kernel to use updated packages.

In []: initiating the repository

```
In [103... import mlflow
import dagshub

# Initialize DagsHub
dagshub.init(repo_owner='nithinyanna3', repo_name='my-first-repo', mlflow=True)
```

Accessing as nithinyanna3

Initialized MLflow to track repo "nithinyanna3/my-first-repo"

Repository nithinyanna3/my-first-repo initialized!

performing logistic regression with pipeline and data transforming

In [159...

```

import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split, StratifiedKFold, GridSearchCV
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler, MinMaxScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import f1_score, confusion_matrix, accuracy_score
import joblib
import mlflow
import mlflow.sklearn
import matplotlib.pyplot as plt

# Assuming secondupdate and test_data are already Loaded DataFrames
# Training data
df = secondupdate.copy() # Use the `secondupdate` DataFrame
X = df.drop("loan_status", axis=1)
y = df["loan_status"]

# Testing data
test_features = test_data.copy()

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_sta

# Separate numerical and categorical columns
numerical_features = X.select_dtypes(include=["int64", "float64"]).columns
categorical_features = X.select_dtypes(include=["object", "category"]).columns

# Preprocessing steps
numerical_transformer = Pipeline(steps=[
    ("scaler", StandardScaler()), # Standardize numerical data
    ("minmax", MinMaxScaler()) # Normalize between 0 and 1
])

categorical_transformer = Pipeline(steps=[
    ("onehot", OneHotEncoder(handle_unknown="ignore")) # One-hot encoding for cate
])

# Combine preprocessors
preprocessor = ColumnTransformer(
    transformers=[
        ("num", numerical_transformer, numerical_features),
        ("cat", categorical_transformer, categorical_features)
    ]
)

# Full pipeline with logistic regression
pipeline = Pipeline(steps=[
    ("preprocessor", preprocessor),
    ("classifier", LogisticRegression(solver="liblinear", random_state=42))
])

# Hyperparameter tuning with GridSearchCV
param_grid = {

```

```

    "classifier__C": [0.01, 0.1, 1, 10, 100], # Regularization strength
    "classifier__penalty": ["l1", "l2"]      # L1/L2 regularization
}

# Stratified K-Fold Cross Validation
cv = StratifiedKFold(n_splits=10, shuffle=True, random_state=42)

# Grid search for best parameters
grid_search = GridSearchCV(estimator=pipe, param_grid=param_grid, cv=cv, scoring='f1')
grid_search.fit(X_train, y_train)

# Best parameters and results
best_model = grid_search.best_estimator_
cv_results = grid_search.cv_results_
mean_f1 = grid_search.best_score_
std_f1 = cv_results["std_test_score"][grid_search.best_index_]

# Predictions and metrics on the test set
y_pred = best_model.predict(X_test)

# Evaluate Accuracy
accuracy = accuracy_score(y_test, y_pred)

# F1-score
f1 = f1_score(y_test, y_pred)

# Confusion Matrix
tn, fp, fn, tp = confusion_matrix(y_test, y_pred).ravel()
print("Best Parameters:", grid_search.best_params_)
print("F1 Score (CV):", mean_f1)
print("Standard Deviation of F1 Score (CV):", std_f1)
print("Accuracy on Test Set:", accuracy)
print("F1 Score on Test Set:", f1)
print("Confusion Matrix (TN, FP, FN, TP):", tn, fp, fn, tp)

# Testing on first 5 rows of test_data
test_predictions = best_model.predict(test_features.head(5))
print("Predictions for the first 5 rows of test_data:", test_predictions)

# Save the model using joblib
joblib.dump(best_model, "logistic_regression_model.pkl")

# Log experiment in MLflow
with mlflow.start_run():
    # Log parameters
    mlflow.log_param("best_C", grid_search.best_params_["classifier__C"])
    mlflow.log_param("best_penalty", grid_search.best_params_["classifier__penalty"])
    mlflow.log_param("cv_splits", 10)

    # Log metrics
    mlflow.log_metric("accuracy", accuracy)
    mlflow.log_metric("f1_score", f1)
    mlflow.log_metric("mean_cv_f1", mean_f1)
    mlflow.log_metric("std_cv_f1", std_f1)
    mlflow.log_metric("TP", tp)
    mlflow.log_metric("TN", tn)

```

```

mlflow.log_metric("FP", fp)
mlflow.log_metric("FN", fn)

# Log model
mlflow.sklearn.log_model(best_model, "logistic_regression_model")

# Save and Log confusion matrix
cm = confusion_matrix(y_test, y_pred)
fig, ax = plt.subplots()
ax.matshow(cm, cmap="coolwarm")
for (i, j), val in np.ndenumerate(cm):
    ax.text(j, i, val, ha="center", va="center")
plt.title("Confusion Matrix")
plt.savefig("confusion_matrix.png")
mlflow.log_artifact("confusion_matrix.png")

print("Experiment successfully logged to MLflow!")

```

Best Parameters: {'classifier__C': 100, 'classifier__penalty': 'l2'}

F1 Score (CV): 0.4728985183873246

Standard Deviation of F1 Score (CV): 0.019636412673607086

Accuracy on Test Set: 0.8857216672407854


F1 Score on Test Set: 0.4736215787385958

Confusion Matrix (TN, FP, FN, TP): 9688 278 1049 597

Predictions for the first 5 rows of test_data: [1 0 1 0 0]

2024/12/17 14:20:58 WARNING mlflow.models.model: Model logged without a signature and input example. Please set `input_example` parameter when logging the model to auto infer the model signature.

Experiment successfully logged to MLflow!



performing ridge classifier with pipeline and data transforming

In [109...

```
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split, StratifiedKFold, GridSearchCV
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.linear_model import RidgeClassifier
from sklearn.metrics import accuracy_score, f1_score, confusion_matrix
import joblib
import mlflow
import mlflow.sklearn
import matplotlib.pyplot as plt

# Assuming `secondupdate` and `test_data` DataFrames are already loaded
# Training data
df = secondupdate.copy() # Use the `secondupdate` DataFrame
X = df.drop("loan_status", axis=1)
y = df["loan_status"]

# Testing data
test_features = test_data.copy()

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_sta

# Separate numerical and categorical columns
numerical_features = X.select_dtypes(include=["int64", "float64"]).columns
categorical_features = X.select_dtypes(include=["object", "category"]).columns

# Preprocessing steps
numerical_transformer = Pipeline(steps=[
    ("scaler", StandardScaler()) # Standardize numerical data
])

categorical_transformer = Pipeline(steps=[
    ("onehot", OneHotEncoder(handle_unknown="ignore")) # One-hot encoding for cate
])

# Combine preprocessors
preprocessor = ColumnTransformer(
    transformers=[
        ("num", numerical_transformer, numerical_features),
        ("cat", categorical_transformer, categorical_features)
    ]
)

# Define the Ridge Classifier model
model = RidgeClassifier()

# Hyperparameter tuning with GridSearchCV
param_grid = {
    'classifier__alpha': [0.1, 1, 10], # Regularization strength
    'classifier__solver': ['auto', 'svd', 'cholesky', 'lsqr', 'sparse_cg'] # Solve
```

```

}

# Full pipeline with model
pipeline = Pipeline(steps=[
    ("preprocessor", preprocessor),
    ("classifier", model)
])

# Stratified K-Fold Cross Validation
cv = StratifiedKFold(n_splits=10, shuffle=True, random_state=42)

# Grid search for best parameters
grid_search = GridSearchCV(estimator=pipeline, param_grid=param_grid, cv=cv, scoring='accuracy')
grid_search.fit(X_train, y_train)

# Best model and parameters
best_model = grid_search.best_estimator_

# Predict on test set
y_pred = best_model.predict(X_test)

# Calculate metrics
accuracy = accuracy_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred, average='weighted')
mean_f1 = np.mean(grid_search.cv_results_["mean_test_score"])
std_f1 = np.std(grid_search.cv_results_["std_test_score"])

# Confusion matrix
cm = confusion_matrix(y_test, y_pred)
tn, fp, fn, tp = cm.ravel()

# Test first 5 rows of `test_data`
test_predictions = best_model.predict(test_features.head(5))
print("Predictions for the first 5 rows of test_data:", test_predictions)

# Save the model using joblib
joblib.dump(best_model, "ridge_classifier_model.pkl")

# Start MLflow Run
with mlflow.start_run():
    # Log parameters
    mlflow.log_param("best_alpha", grid_search.best_params_["classifier__alpha"])
    mlflow.log_param("best_solver", grid_search.best_params_["classifier__solver"])
    mlflow.log_param("cv_splits", 10)

    # Log metrics
    mlflow.log_metric("accuracy", accuracy)
    mlflow.log_metric("f1_score", f1)
    mlflow.log_metric("mean_cv_f1", mean_f1)
    mlflow.log_metric("std_cv_f1", std_f1)
    mlflow.log_metric("TP", tp)
    mlflow.log_metric("TN", tn)
    mlflow.log_metric("FP", fp)
    mlflow.log_metric("FN", fn)

    # Log model

```



```

mlflow.sklearn.log_model(best_model, "ridge_classifier_model")

# Save and Log confusion matrix
fig, ax = plt.subplots()
ax.matshow(cm, cmap="coolwarm")
for (i, j), val in np.ndenumerate(cm):
    ax.text(j, i, val, ha="center", va="center")
plt.title("Confusion Matrix")
plt.savefig("confusion_matrix.png")
mlflow.log_artifact("confusion_matrix.png")

# Print metrics to notebook
print("Accuracy:", accuracy)
print("F1 Score (Weighted):", f1)
print("Mean F1 Score (CV):", mean_f1)
print("Std F1 Score (CV):", std_f1)
print("Confusion Matrix:")
print(cm)
print("True Positives (TP):", tp)
print("False Negatives (FN):", fn)
print("False Positives (FP):", fp)
print("True Negatives (TN):", tn)

print("Experiment successfully logged to MLflow!")

```

Predictions for the first 5 rows of test_data: [0 0 0 0 0]

2024/12/17 13:22:08 WARNING mlflow.models.model: Model logged without a signature and input example. Please set `input_example` parameter when logging the model to auto infer the model signature.

Accuracy: 0.880985187736824

F1 Score (Weighted): 0.8508153925100107

Mean F1 Score (CV): 0.8808867768311529

Std F1 Score (CV): 1.6995211594251085e-05

Confusion Matrix:

```
[[9868  98]
 [1284 362]]
```


True Positives (TP): 362


False Negatives (FN): 1284


False Positives (FP): 98

True Negatives (TN): 9868

Experiment successfully logged to MLflow!

 View run mysterious-goat-182 at: <https://dagshub.com/nithinyanna3/my-first-repo.mlflow/#/experiments/0/runs/ab78a0a385734998a0e00c5c5ba607f1>

 View experiment at: <https://dagshub.com/nithinyanna3/my-first-repo.mlflow/#/experiments/0>



performing XGBClassifier with pipeline and data transforming

```
In [113... import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split, StratifiedKFold, GridSearchCV
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from xgboost import XGBClassifier
from sklearn.metrics import accuracy_score, f1_score, confusion_matrix
import mlflow
import mlflow.sklearn
import matplotlib.pyplot as plt
import joblib

# Assuming `secondupdate` and `test_data` DataFrames are already loaded
# Training data
df = secondupdate.copy() # Use the `secondupdate` DataFrame
X = df.drop("loan_status", axis=1)
y = df["loan_status"]

# Testing data
test_features = test_data.copy()

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_sta

# Separate numerical and categorical columns
numerical_features = X.select_dtypes(include=["int64", "float64"]).columns
```

```

categorical_features = X.select_dtypes(include=["object", "category"]).columns

# Preprocessing steps
numerical_transformer = Pipeline(steps=[
    ("scaler", StandardScaler()) # Standardize numerical data
])

categorical_transformer = Pipeline(steps=[
    ("onehot", OneHotEncoder(handle_unknown="ignore")) # One-hot encoding for cate
])

# Combine preprocessors
preprocessor = ColumnTransformer(
    transformers=[
        ("num", numerical_transformer, numerical_features),
        ("cat", categorical_transformer, categorical_features)
    ]
)

# Define the XGB Classifier model
model = XGBClassifier(random_state=42, use_label_encoder=False, eval_metric='mloglo

# Hyperparameter tuning with GridSearchCV
param_grid = {
    'classifier__n_estimators': [50, 100, 200], # Number of trees in the forest
    'classifier__max_depth': [3, 6, 10], # Maximum depth of the tree
    'classifier__learning_rate': [0.01, 0.1, 0.2], # Learning rate
    'classifier__subsample': [0.8, 1.0], # Fraction of samples used for fitting
    'classifier__colsample_bytree': [0.8, 1.0] # Fraction of features used for eac
}

# Full pipeline with model
pipeline = Pipeline(steps=[
    ("preprocessor", preprocessor),
    ("classifier", model)
])

# Stratified K-Fold Cross Validation
cv = StratifiedKFold(n_splits=10, shuffle=True, random_state=42)

# Grid search for best parameters
import time

# Timer starts
start_time = time.time()

print("Starting GridSearchCV...")
grid_search = GridSearchCV(
    estimator=pipeline,
    param_grid=param_grid,
    cv=cv,
    scoring="accuracy",
    n_jobs=-1,
    verbose=2 # Verbosity level to show progress
)

```

```
grid_search.fit(X_train, y_train)

# Timer ends
end_time = time.time()
execution_time = end_time - start_time

print(f"\nGridSearchCV completed in {execution_time:.2f} seconds.")

# Best model and parameters
best_model = grid_search.best_estimator_

# Predict on test set
y_pred = best_model.predict(X_test)

# Calculate metrics
accuracy = accuracy_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred, average='weighted')
mean_f1 = np.mean(grid_search.cv_results_["mean_test_score"])
std_f1 = np.std(grid_search.cv_results_["std_test_score"])

# Confusion matrix
cm = confusion_matrix(y_test, y_pred)
tn, fp, fn, tp = cm.ravel()

# Test first 5 rows of `test_data`
test_predictions = best_model.predict(test_features.head(5))
print("\nPredictions for the first 5 rows of test_data:", test_predictions)

# Save the model using joblib
joblib.dump(best_model, "xgb_classifier_model.pkl")
print("Model saved as 'xgb_classifier_model.pkl'.")

# Start MLflow Run
with mlflow.start_run():
    # Log parameters
    mlflow.log_param("best_n_estimators", grid_search.best_params_["classifier__n_estimators"])
    mlflow.log_param("best_max_depth", grid_search.best_params_["classifier__max_depth"])
    mlflow.log_param("best_learning_rate", grid_search.best_params_["classifier__learning_rate"])
    mlflow.log_param("best_subsample", grid_search.best_params_["classifier__subsample"])
    mlflow.log_param("best_colsample_bytree", grid_search.best_params_["classifier__colsample_bytree"])
    mlflow.log_param("cv_splits", 10)

    # Log metrics
    mlflow.log_metric("accuracy", accuracy)
    mlflow.log_metric("f1_score", f1)
    mlflow.log_metric("mean_cv_f1", mean_f1)
    mlflow.log_metric("std_cv_f1", std_f1)
    mlflow.log_metric("TP", tp)
    mlflow.log_metric("TN", tn)
    mlflow.log_metric("FP", fp)
    mlflow.log_metric("FN", fn)

    # Log model
    mlflow.sklearn.log_model(best_model, "xgb_classifier_model")

    # Save and Log confusion matrix
```

```

fig, ax = plt.subplots()
ax.matshow(cm, cmap="coolwarm")
for (i, j), val in np.ndenumerate(cm):
    ax.text(j, i, val, ha="center", va="center")
plt.title("Confusion Matrix")
plt.savefig("confusion_matrix.png")
mlflow.log_artifact("confusion_matrix.png")

# Print metrics to notebook
print("\nMetrics:")
print("Accuracy:", accuracy)
print("F1 Score (Weighted):", f1)
print("Mean F1 Score (CV):", mean_f1)
print("Std F1 Score (CV):", std_f1)
print("Confusion Matrix:")
print(cm)
print("True Positives (TP):", tp)
print("False Negatives (FN):", fn)
print("False Positives (FP):", fp)
print("True Negatives (TN):", tn)

print("Experiment successfully logged to MLflow!")

```

Starting GridSearchCV...

Fitting 10 folds for each of 108 candidates, totalling 1080 fits

C:\Users\HP\anaconda3\Lib\site-packages\xgboost\core.py:158: UserWarning: [13:31:42]
 WARNING: C:\buildkite-agent\builds\buildkite-windows-cpu-autoscaling-group-i-0c55ff5
 f71b100e98-1\xgboost\xgboost-ci-windows\src\learner.cc:740:
 Parameters: { "use_label_encoder" } are not used.

warnings.warn(msg, UserWarning)

GridSearchCV completed in 161.41 seconds.

Predictions for the first 5 rows of test_data: [1 0 1 0 0]

Model saved as 'xgb_classifier_model.pkl'.

2024/12/17 13:31:53 WARNING mlflow.models.model: Model logged without a signature and input example. Please set `input_example` parameter when logging the model to auto infer the model signature.

Metrics:

Accuracy: 0.9173269032035825

F1 Score (Weighted): 0.9139787570137752

Mean F1 Score (CV): 0.906618580447282

Std F1 Score (CV): 0.0016429811984086614

Confusion Matrix:

[[9629 337]

[623 1023]]


True Positives (TP): 1023


False Negatives (FN): 623


False Positives (FP): 337

True Negatives (TN): 9629

Experiment successfully logged to MLflow!

 View run marvelous-crab-821 at: <https://dagshub.com/nithinyanna3/my-first-repo.mlflow/#/experiments/0/runs/07275b51f40a43719b58c994bef13fd2>

 View experiment at: <https://dagshub.com/nithinyanna3/my-first-repo.mlflow/#/experiments/0>



performing Random forest with pipeline and data transforming

```
In [116... import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split, StratifiedKFold, GridSearchCV
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, f1_score, confusion_matrix
import mlflow
import mlflow.sklearn
import matplotlib.pyplot as plt
import joblib

# Assuming `secondupdate` and `test_data` DataFrames are already loaded
# Training data
df = secondupdate.copy() # Use the `secondupdate` DataFrame
X = df.drop("loan_status", axis=1)
y = df["loan_status"]

# Testing data
test_features = test_data.copy()

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_sta

# Separate numerical and categorical columns
numerical_features = X.select_dtypes(include=["int64", "float64"]).columns
```

```

categorical_features = X.select_dtypes(include=["object", "category"]).columns

# Preprocessing steps
numerical_transformer = Pipeline(steps=[
    ("scaler", StandardScaler()) # Standardize numerical data
])

categorical_transformer = Pipeline(steps=[
    ("onehot", OneHotEncoder(handle_unknown="ignore")) # One-hot encoding for cate
])

# Combine preprocessors
preprocessor = ColumnTransformer(
    transformers=[
        ("num", numerical_transformer, numerical_features),
        ("cat", categorical_transformer, categorical_features)
    ]
)

# Define the Random Forest Classifier model
model = RandomForestClassifier(random_state=42)

# Hyperparameter tuning with GridSearchCV
param_grid = {
    'classifier__n_estimators': [50, 100, 200], # Number of trees in the forest
    'classifier__max_depth': [None, 10, 20], # Maximum depth of the tree
    'classifier__min_samples_split': [2, 5, 10], # Minimum number of samples requi
    'classifier__min_samples_leaf': [1, 2, 4] # Minimum number of samples required
}

# Full pipeline with model
pipeline = Pipeline(steps=[
    ("preprocessor", preprocessor),
    ("classifier", model)
])

# Stratified K-Fold Cross Validation
cv = StratifiedKFold(n_splits=10, shuffle=True, random_state=42)

# Grid search for best parameters
print("Starting GridSearchCV...")
grid_search = GridSearchCV(estimator=pipeline, param_grid=param_grid, cv=cv, scorin
grid_search.fit(X_train, y_train)

# Best model and parameters
best_model = grid_search.best_estimator_

# Predict on test set
y_pred = best_model.predict(X_test)

# Calculate metrics
accuracy = accuracy_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred, average='weighted')
mean_f1 = np.mean(grid_search.cv_results_["mean_test_score"])
std_f1 = np.std(grid_search.cv_results_["std_test_score"])

```

```

# Confusion matrix
cm = confusion_matrix(y_test, y_pred)
tn, fp, fn, tp = cm.ravel()

# Test first 5 rows of `test_data`
test_predictions = best_model.predict(test_features.head(5))
print("\nPredictions for the first 5 rows of test_data:", test_predictions)

# Save the model using joblib
joblib.dump(best_model, "random_forest_classifier_model.pkl")
print("Model saved as 'random_forest_classifier_model.pkl'.")

# Start MLflow Run
with mlflow.start_run():
    # Log parameters
    mlflow.log_param("best_n_estimators", grid_search.best_params_["classifier__n_estimators"])
    mlflow.log_param("best_max_depth", grid_search.best_params_["classifier__max_depth"])
    mlflow.log_param("best_min_samples_split", grid_search.best_params_["classifier__min_samples_split"])
    mlflow.log_param("best_min_samples_leaf", grid_search.best_params_["classifier__min_samples_leaf"])
    mlflow.log_param("cv_splits", 10)

    # Log metrics
    mlflow.log_metric("accuracy", accuracy)
    mlflow.log_metric("f1_score", f1)
    mlflow.log_metric("mean_cv_f1", mean_f1)
    mlflow.log_metric("std_cv_f1", std_f1)
    mlflow.log_metric("TP", tp)
    mlflow.log_metric("TN", tn)
    mlflow.log_metric("FP", fp)
    mlflow.log_metric("FN", fn)

    # Log model
    mlflow.sklearn.log_model(best_model, "random_forest_classifier_model")

    # Save and Log confusion matrix
    fig, ax = plt.subplots()
    ax.matshow(cm, cmap="coolwarm")
    for (i, j), val in np.ndenumerate(cm):
        ax.text(j, i, val, ha="center", va="center")
    plt.title("Confusion Matrix")
    plt.savefig("confusion_matrix.png")
    mlflow.log_artifact("confusion_matrix.png")

    # Print metrics to notebook
    print("\nMetrics:")
    print("Accuracy:", accuracy)
    print("F1 Score (Weighted):", f1)
    print("Mean F1 Score (CV):", mean_f1)
    print("Std F1 Score (CV):", std_f1)
    print("Confusion Matrix:")
    print(cm)
    print("True Positives (TP):", tp)
    print("False Negatives (FN):", fn)
    print("False Positives (FP):", fp)
    print("True Negatives (TN):", tn)

```



```
print("Experiment successfully logged to MLflow!")
```

Starting GridSearchCV...

Fitting 10 folds for each of 81 candidates, totalling 810 fits

Predictions for the first 5 rows of test_data: [1 0 1 0 0]

Model saved as 'random_forest_classifier_model.pkl'.

2024/12/17 13:49:18 WARNING mlflow.models.model: Model logged without a signature and input example. Please set `input_example` parameter when logging the model to auto infer the model signature.

Metrics:

Accuracy: 0.9099207716155701

F1 Score (Weighted): 0.9056951580786216

Mean F1 Score (CV): 0.9126145944244894

Std F1 Score (CV): 0.0003477845852202995

Confusion Matrix:

```
[[9606  360]
```

```
 [ 686  960]]
```


True Positives (TP): 960


False Negatives (FN): 686


False Positives (FP): 360

True Negatives (TN): 9606

Experiment successfully logged to MLflow!

 View run flawless-midge-434 at: <https://dagshub.com/nithinyanna3/my-first-repo.mlflow/#/experiments/0/runs/e68d2fa0504f4a009868e8353f4dd6d7>

 View experiment at: <https://dagshub.com/nithinyanna3/my-first-repo.mlflow/#/experiments/0>



```
In [151... import mlflow
import dagshub

# Initialize DagsHub
dagshub.init(repo_owner='nithinyanna3', repo_name='my-first-repo', mlflow=True)
```

Initialized MLflow to track repo "nithinyanna3/my-first-repo"

Repository nithinyanna3/my-first-repo initialized!

applying column transformations and performing all the 4 models and logging into dagshub

```
In [184... import numpy as np
import pandas as pd
import tempfile
import joblib
from sklearn.model_selection import train_test_split
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.linear_model import LogisticRegression, RidgeClassifier
from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier
from sklearn.metrics import accuracy_score, f1_score, confusion_matrix
import mlflow
import mlflow.sklearn
import matplotlib.pyplot as plt
import seaborn as sns

# Feature Engineering Function
def feature_engineering(df):
    df = df.copy()

    # Derived features
    df['loan_to_income_ratio'] = df['loan_amnt'] / (df['person_income'] + 1e-5)
    df['income_to_age_ratio'] = df['person_income'] / (df['person_age'] + 1e-5)
    df['credit_age_ratio'] = df['cb_person_cred_hist_length'] / (df['person_age'] + 1e-5)
    df['high_loan'] = (df['loan_amnt'] > 10000).astype(int)
    df['log_loan_amnt'] = np.log1p(df['loan_amnt'])
    df['log_income'] = np.log1p(df['person_income'])

    # Return modified DataFrame and List of new features
    new_features = [
        'loan_to_income_ratio', 'income_to_age_ratio', 'credit_age_ratio',
        'high_loan', 'log_loan_amnt', 'log_income'
    ]
    return df, new_features

# Process the training data (secondupdate)
print("Processing training data...")
thirdupdate, new_features = feature_engineering(secondupdate)
print("Training data processed. Shape:", thirdupdate.shape)

# Process the testing data (test_data)
print("Processing test data...")
test_data_fe, _ = feature_engineering(test_data)
```

```

print("Test data processed. Shape:", test_data_fe.shape)

# Save feature-engineered datasets (optional)
thirdupdate.to_csv("thirdupdate.csv", index=False)
test_data_fe.to_csv("test_data_fe.csv", index=False)

# Prepare training and target variables
X = thirdupdate.drop("loan_status", axis=1) # Independent variables
y = thirdupdate["loan_status"] # Target variable

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_sta

# Identify numerical and categorical columns
numerical_features = X_train.select_dtypes(include=["int64", "float64"]).columns
categorical_features = X_train.select_dtypes(include=["object", "category"]).column

# Preprocessing Pipelines
numerical_transformer = Pipeline(steps=[
    ("scaler", StandardScaler())
])

categorical_transformer = Pipeline(steps=[
    ("onehot", OneHotEncoder(handle_unknown="ignore"))
])

preprocessor = ColumnTransformer(transformers=[
    ("num", numerical_transformer, numerical_features),
    ("cat", categorical_transformer, categorical_features)
])

# Define models
models = {
    "LogisticRegression": LogisticRegression(max_iter=500, random_state=42),
    "RidgeClassifier": RidgeClassifier(random_state=42),
    "RandomForestClassifier": RandomForestClassifier(random_state=42),
    "XGBClassifier": XGBClassifier(random_state=42, use_label_encoder=False, eval_m
}

# Set the MLflow tracking URI and experiment name
mlflow.set_tracking_uri("https://dagshub.com/nithinyanna3/my-first-repo.mlflow") #
mlflow.set_experiment("Default") # Set experiment name

# Train and Log Models
for model_name, model in models.items():
    with mlflow.start_run(run_name=f"{model_name}_FeatureEngineering"):
        print(f"\nTraining {model_name}...")

        # Full pipeline
        pipeline = Pipeline(steps=[
            ("preprocessor", preprocessor),
            ("classifier", model)
        ])

        # Train the model
        pipeline.fit(X_train, y_train)

```

```
# Predict
y_pred = pipeline.predict(X_test)
acc = accuracy_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred, average='weighted')
cm = confusion_matrix(y_test, y_pred)

# Print metrics
print(f"Model: {model_name}")
print(f"Accuracy: {acc:.4f}")
print(f"F1 Score: {f1:.4f}")
print(f"Confusion Matrix:\n{cm}")

# Extract confusion matrix values
tp, fn, fp, tn = cm.ravel()

# Print confusion matrix values
print(f"True Positives (TP): {tp}")
print(f"True Negatives (TN): {tn}")
print(f"False Positives (FP): {fp}")
print(f"False Negatives (FN): {fn}")

# Visualize confusion matrix
plt.figure(figsize=(6, 4))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", xticklabels=np.unique(y))
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title(f'Confusion Matrix - {model_name}')
plt.show()

# Save the model using joblib
model_filename = f"{model_name}_pipeline.joblib"
joblib.dump(pipeline, model_filename)
print(f"Model saved as {model_filename}")

# Log metrics and model to MLflow
mlflow.log_param("model_name", model_name)
mlflow.log_metric("accuracy", acc)
mlflow.log_metric("f1_score", f1)
mlflow.log_metric("TP", tp)
mlflow.log_metric("TN", tn)
mlflow.log_metric("FP", fp)
mlflow.log_metric("FN", fn)
mlflow.sklearn.log_model(pipeline, f"{model_name}_model")

# Predict for the first 5 rows of test data
print(f"\nPredictions on first 5 rows of test_data_fe using {model_name}:")
predictions = pipeline.predict(test_data_fe.head(5))
print(predictions)

print("\nAll models have been trained, logged, and saved successfully.")
```

```
Processing training data...
Training data processed. Shape: (58058, 16)
Processing test data...
Test data processed. Shape: (39098, 15)
```

```
Training LogisticRegression...
```

```
Model: LogisticRegression
```

```
Accuracy: 0.8880
```

```
F1 Score: 0.8731
```

```
Confusion Matrix:
```

```
[[9697 269]
```


```
 [1032 614]]
```

```
True Positives (TP): 9697
```

```
True Negatives (TN): 614
```

```
False Positives (FP): 1032
```

```
False Negatives (FN): 269
```





```
Model saved as LogisticRegression_pipeline.joblib
```

```
2024/12/17 14:29:30 WARNING mlflow.models.model: Model logged without a signature and input example. Please set `input_example` parameter when logging the model to auto infer the model signature.
```

Predictions on first 5 rows of test_data_fe using LogisticRegression:

```
[1 0 1 0 0]
```

 View run LogisticRegression_FeatureEngineering at: <https://dagshub.com/nithinyanna3/my-first-repo.mlflow/#/experiments/0/runs/8b49ef5c002744b4b5664f42119eddf3>

 View experiment at: <https://dagshub.com/nithinyanna3/my-first-repo.mlflow/#/experiments/0>

Training RidgeClassifier...

Model: RidgeClassifier

Accuracy: 0.8826

F1 Score: 0.8563

Confusion Matrix:

```
[[9834 132]
```


```
 [1231 415]]
```

True Positives (TP): 9834

True Negatives (TN): 415

False Positives (FP): 1231

False Negatives (FN): 132





Model saved as RidgeClassifier_pipeline.joblib

2024/12/17 14:29:42 WARNING mlflow.models.model: Model logged without a signature and input example. Please set `input_example` parameter when logging the model to auto infer the model signature.

Predictions on first 5 rows of test_data_fe using RidgeClassifier:

```
[1 0 0 0 0]
```

 View run RidgeClassifier_FeatureEngineering at: <https://dagshub.com/nithinyanna3/my-first-repo.mlflow/#/experiments/0/runs/a065305fdb5d44968434ebe6c4e56f9d>

 View experiment at: <https://dagshub.com/nithinyanna3/my-first-repo.mlflow/#/experiments/0>

Training RandomForestClassifier...

Model: RandomForestClassifier

Accuracy: 0.9114

F1 Score: 0.9073

Confusion Matrix:

```
[[9613  353]
```


```
 [ 676  970]]
```

True Positives (TP): 9613

True Negatives (TN): 970

False Positives (FP): 676

False Negatives (FN): 353





Model saved as RandomForestClassifier_pipeline.joblib

2024/12/17 14:30:00 WARNING mlflow.models.model: Model logged without a signature and input example. Please set `input_example` parameter when logging the model to auto infer the model signature.

Predictions on first 5 rows of test_data_fe using RandomForestClassifier:

```
[1 0 1 0 0]
```

 View run RandomF-Fc at: <https://dagshub.com/nithinyanna3/my-first-repo.mlflow/#/experiments/0/runs/3b554e23bfd540e0a3a90200bf75c977>

 View experiment at: <https://dagshub.com/nithinyanna3/my-first-repo.mlflow/#/experiments/0>

Training XGBClassifier...

```
C:\Users\HP\anaconda3\Lib\site-packages\xgboost\core.py:158: UserWarning: [14:30:32]
WARNING: C:\buildkite-agent\builds\buildkite-windows-cpu-autoscaling-group-i-0c55ff5
f71b100e98-1\xgboost\xgboost-ci-windows\src\learner.cc:740:
Parameters: { "use_label_encoder" } are not used.
```

```
warnings.warn(smsg, UserWarning)
```

Model: XGBClassifier

Accuracy: 0.9179

F1 Score: 0.9150

Confusion Matrix:


```
[[9618  348]
 [ 605 1041]]
```

True Positives (TP): 9618

True Negatives (TN): 1041

False Positives (FP): 605

False Negatives (FN): 348



Model saved as XGBClassifier_pipeline.joblib

```
2024/12/17 14:30:38 WARNING mlflow.models.model: Model logged without a signature and
input example. Please set `input_example` parameter when logging the model to auto
infer the model signature.
```

Predictions on first 5 rows of test_data_fe using XGBClassifier:

```
[1 0 1 0 0]
```

View run XGBClassifier_FeatureEngineering at: <https://dagshub.com/nithinyanna3/my-first-repo.mlflow/#/experiments/0/runs/f7171e6b1e654dd3b58dcd1aedbd806a>

View experiment at: <https://dagshub.com/nithinyanna3/my-first-repo.mlflow/#/experiments/0>

All models have been trained, logged, and saved successfully.

applying variance threshold for the dataset and refining the dataset

In [196...

```

import pandas as pd
from sklearn.feature_selection import VarianceThreshold

# Assuming dfsecondupdate is already loaded in the environment
# dfsecondupdate = pd.read_csv("path_to_your_dfsecondupdate.csv") # Uncomment and

# Create a copy of dfsecondupdate
df = secondupdate.copy()

# Separate target variable
X = df.drop("loan_status", axis=1) # Features (excluding target)
y = df["loan_status"] # Target variable

# Step 1: Apply Variance Threshold to remove Low variance features
variance_threshold = VarianceThreshold(threshold=0.01) # You can adjust the thresh

# Apply VarianceThreshold to the features (X)
X_var = variance_threshold.fit_transform(X)

# Step 2: Identify the columns that were removed
removed_columns = X.columns[~variance_threshold.get_support()]

# Step 3: Convert the resulting array back to a DataFrame
X_var_df = pd.DataFrame(X_var, columns=X.columns[variance_threshold.get_support()])

# Add the target variable back to the DataFrame
X_var_df["loan_status"] = y

# Step 4: Save the dataset with Low-variance features removed
output_path = "fourthupdate.csv" # Save in the environment
X_var_df.to_csv(output_path, index=False)

# Print details
print(f"Low-variance features removed: {' '.join(removed_columns)}")
print(f"Original dataset size: {X.shape}")
print(f"New dataset size after removing low-variance features: {X_var_df.shape}")

# Print the columns in secondupdate and fourthupdate to see the difference
print(f"Columns in secondupdate: {X.columns.tolist()}")
print(f"Columns in fourthupdate: {X_var_df.columns.tolist()}")

```

Low-variance features removed: loan_grade, loan_percent_income, cb_person_default_on_file

Original dataset size: (58058, 9)

New dataset size after removing low-variance features: (58058, 7)

Columns in secondupdate: ['person_age', 'person_income', 'person_emp_length', 'loan_grade', 'loan_amnt', 'loan_int_rate', 'loan_percent_income', 'cb_person_default_on_file', 'cb_person_cred_hist_length']

Columns in fourthupdate: ['person_age', 'person_income', 'person_emp_length', 'loan_amnt', 'loan_int_rate', 'cb_person_cred_hist_length', 'loan_status']

knowing the highly correlated variables

In [212...

```

import pandas as pd
import seaborn as sns

```

```
import matplotlib.pyplot as plt

# Load the dataset after applying Variance Threshold
df_variance_threshold = pd.read_csv("fourthupdate.csv")

# Separate target variable
X_var = df_variance_threshold.drop("loan_status", axis=1) # Features (excluding ta
y = df_variance_threshold["loan_status"] # Target variable

# Calculate the correlation matrix for the features (X_var)
correlation_matrix = X_var.corr()

# Print the correlation matrix
print("Correlation Matrix:")
print(correlation_matrix)

# Extract high correlated variables with a threshold, e.g., |correlation| > 0.7
correlation_threshold = 0.7
highly_correlated_vars = correlation_matrix[abs(correlation_matrix) > correlation_t
highly_correlated_vars = highly_correlated_vars[highly_correlated_vars != 1] # Exc

# Print highly correlated variables
print(f"High correlated variables (threshold > {correlation_threshold}):")
print(highly_correlated_vars)

# Plotting the heatmap for better visualization
plt.figure(figsize=(12, 8))
sns.heatmap(correlation_matrix, annot=True, cmap="coolwarm", fmt=".2f", linewidths=
plt.title("Correlation Matrix Heatmap")
plt.show()
```

Correlation Matrix:

	person_age	person_income	person_emp_length	\
person_age	1.000000	0.115944	0.116679	
person_income	0.115944	1.000000	0.199807	
person_emp_length	0.116679	0.199807	1.000000	
loan_amnt	0.047632	0.378064	0.092139	
loan_int_rate	0.005217	-0.074235	-0.108780	
cb_person_cred_hist_length	0.871600	0.095027	0.101096	


	loan_amnt	loan_int_rate	\
person_age	0.047632	0.005217	
person_income	0.378064	-0.074235	
person_emp_length	0.092139	-0.108780	
loan_amnt	1.000000	0.113635	
loan_int_rate	0.113635	1.000000	
cb_person_cred_hist_length	0.041843	0.003802	

	cb_person_cred_hist_length
person_age	0.871600
person_income	0.095027
person_emp_length	0.101096
loan_amnt	0.041843
loan_int_rate	0.003802
cb_person_cred_hist_length	1.000000

High correlated variables (threshold > 0.7):

person_age	cb_person_cred_hist_length	0.8716
cb_person_cred_hist_length	person_age	0.8716

dtype: float64



In [206...

```
import pandas as pd

# Assuming the file is saved in the environment as 'updatedtestdata.csv'
file_path = "updatedtestdata.csv"

# Load the dataset
df_updated_test = pd.read_csv(file_path)

# Display the first few rows of the dataset
print(df_updated_test.head())
```

	person_age	person_income	person_emp_length	loan_grade	loan_amnt	\
0	23	69000	3.0	5.0	25000	
1	26	96000	6.0	2.0	10000	
2	26	30000	5.0	4.0	4000	
3	33	50000	4.0	0.0	7000	
4	26	102000	8.0	3.0	15000	

	loan_int_rate	loan_percent_income	cb_person_default_on_file	\
0	15.76	0.36	1	
1	12.68	0.10	2	
2	17.19	0.13	2	
3	8.90	0.14	1	
4	16.32	0.15	2	

	cb_person_cred_hist_length
0	2
1	4
2	2
3	7
4	4

In []: updating the test data accordingly

```
In [209... import pandas as pd

# Assuming the file is saved in the environment as 'updatedtestdata.csv'
file_path = "updatedtestdata.csv"

# Load the dataset
df_updated_test = pd.read_csv(file_path)

# Columns to be removed (replace with your desired columns)
columns_to_remove = ['loan_grade', 'cb_person_default_on_file', 'loan_percent_income']

# Remove specified columns
df_test_variance = df_updated_test.drop(columns=columns_to_remove)

# Save the new dataframe as 'test_variance.csv'
output_path = "test_variance.csv"
df_test_variance.to_csv(output_path, index=False)

# Display the first few rows of the modified dataset
print(df_test_variance.head())
```

	person_age	person_income	person_emp_length	loan_amnt	loan_int_rate	\
0	23	69000	3.0	25000	15.76	
1	26	96000	6.0	10000	12.68	
2	26	30000	5.0	4000	17.19	
3	33	50000	4.0	7000	8.90	
4	26	102000	8.0	15000	16.32	

	cb_person_cred_hist_length
0	2
1	4
2	2
3	7
4	4

deplyomnet along with logging to dagshub

```
In [228... import pandas as pd
import numpy as np
import joblib
from sklearn.model_selection import train_test_split
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.linear_model import LogisticRegression, RidgeClassifier
from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier
from sklearn.metrics import accuracy_score, f1_score, confusion_matrix
import mlflow
import mlflow.sklearn
import matplotlib.pyplot as plt
import seaborn as sns

# Load data from fourthupdate.csv
df = pd.read_csv("fourthupdate.csv")

# Separate target variable
X = df.drop("loan_status", axis=1) # Features (excluding target)
y = df["loan_status"] # Target variable

# Handling NaN values
X.dropna(inplace=True)
y.dropna(inplace=True)

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_sta

# Identify numerical and categorical columns
numerical_features = X_train.select_dtypes(include=["int64", "float64"]).columns
categorical_features = X_train.select_dtypes(include=["object", "category"]).column

# Preprocessing Pipelines
numerical_transformer = Pipeline(steps=[
    ("scaler", StandardScaler())
])

categorical_transformer = Pipeline(steps=[
```

```

    ("onehot", OneHotEncoder(handle_unknown="ignore"))
])

preprocessor = ColumnTransformer(transformers=[
    ("num", numerical_transformer, numerical_features),
    ("cat", categorical_transformer, categorical_features)
])

mlflow.set_tracking_uri("https://dagshub.com/nithinyanna3/my-first-repo.mlflow") #
mlflow.set_experiment("Default")

# Define models
models = {
    "LogisticRegression": LogisticRegression(max_iter=500, random_state=42),
    "RidgeClassifier": RidgeClassifier(random_state=42),
    "RandomForestClassifier": RandomForestClassifier(random_state=42),
    "XGBClassifier": XGBClassifier(random_state=42, use_label_encoder=False, eval_m
}

# Training and Logging Models
for model_name, model in models.items():
    with mlflow.start_run(run_name=f"{model_name}_Variance"):
        print(f"\nTraining {model_name}...")

        # Full pipeline
        pipeline = Pipeline(steps=[
            ("preprocessor", preprocessor),
            ("classifier", model)
        ])

        # Train the model
        pipeline.fit(X_train, y_train)

        # Predict
        y_pred = pipeline.predict(X_test)
        acc = accuracy_score(y_test, y_pred)
        f1 = f1_score(y_test, y_pred, average='weighted')
        cm = confusion_matrix(y_test, y_pred)

        # Print metrics
        print(f"Model: {model_name}")
        print(f"Accuracy: {acc:.4f}")
        print(f"F1 Score: {f1:.4f}")
        print(f"Confusion Matrix:\n{cm}")

        # Extract confusion matrix values
        tp, fn, fp, tn = cm.ravel()

        # Print confusion matrix values
        print(f"True Positives (TP): {tp}")
        print(f"True Negatives (TN): {tn}")
        print(f"False Positives (FP): {fp}")
        print(f"False Negatives (FN): {fn}")

        # Visualize confusion matrix
        plt.figure(figsize=(6, 4))

```

```
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", xticklabels=np.unique(y)
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title(f'Confusion Matrix - {model_name}')
plt.show()

# Save the model using joblib
model_filename = f"{model_name}_pipeline.joblib"
joblib.dump(pipeline, model_filename)
print(f"Model saved as {model_filename}")

# Log metrics and model to MLflow
mlflow.log_param("model_name", model_name)
mlflow.log_metric("accuracy", acc)
mlflow.log_metric("f1_score", f1)
mlflow.log_metric("TP", tp)
mlflow.log_metric("TN", tn)
mlflow.log_metric("FP", fp)
mlflow.log_metric("FN", fn)

# Predict for the first 5 rows of test data
print(f"\nPredictions on first 5 rows of test_variance using {model_name}:")
test_variance = pd.read_csv("test_variance.csv")
predictions = pipeline.predict(test_variance.head(5))
print(predictions)

print("\nAll models have been trained, logged, and saved successfully.")
```

Training LogisticRegression...

Model: LogisticRegression

Accuracy: 0.8582

F1 Score: 0.7927

Confusion Matrix:


```
[[9865   0]
 [1630   0]]
```

True Positives (TP): 9865

True Negatives (TN): 0

False Positives (FP): 1630


False Negatives (FN): 0



Model saved as LogisticRegression_pipeline.joblib

Predictions on first 5 rows of test_variance using LogisticRegression:

[0. 0. 0. 0. 0.]

 View run LogisticRegression_Variance at: <https://dagshub.com/nithinyanna3/my-first-repo.mlflow/#/experiments/0/runs/7adc4544b4e048f9ae6f148345ce147c>

 View experiment at: <https://dagshub.com/nithinyanna3/my-first-repo.mlflow/#/experiments/0>

Training RidgeClassifier...

Model: RidgeClassifier

Accuracy: 0.8582

F1 Score: 0.7927

Confusion Matrix:

[[9865 0]


[1630 0]]

True Positives (TP): 9865

True Negatives (TN): 0

False Positives (FP): 1630

False Negatives (FN): 0




Model saved as RidgeClassifier_pipeline.joblib

Predictions on first 5 rows of test_variance using RidgeClassifier:

[0. 0. 0. 0. 0.]

 View run RidgeClassifier_Variance at: <https://dagshub.com/nithinyanna3/my-first-repo.mlflow/#/experiments/0/runs/09fcc6bef97448b190f27e428ce5a0bf>

 View experiment at: <https://dagshub.com/nithinyanna3/my-first-repo.mlflow/#/experiments/0>

Training RandomForestClassifier...

Model: RandomForestClassifier

Accuracy: 0.8554

F1 Score: 0.7923

Confusion Matrix:

[[9827 38]


[1624 6]]

True Positives (TP): 9827

True Negatives (TN): 6

False Positives (FP): 1624


False Negatives (FN): 38




Model saved as RandomForestClassifier_pipeline.joblib

Predictions on first 5 rows of test_variance using RandomForestClassifier:

[0. 0. 0. 0. 0.]

 View run RandomForestClassifier_Variance at: <https://dagshub.com/nithinyanna3/my-first-repo.mlflow/#/experiments/0/runs/52644c1f354c4987b26ee160b363833d>

 View experiment at: <https://dagshub.com/nithinyanna3/my-first-repo.mlflow/#/experiments/0>

Training XGBClassifier...

Model: XGBClassifier

Accuracy: 0.8574

F1 Score: 0.7927

Confusion Matrix:

[[9854 11]

[1628 2]]

True Positives (TP): 9854


True Negatives (TN): 2

False Positives (FP): 1628

False Negatives (FN): 11

C:\Users\HP\anaconda3\Lib\site-packages\xgboost\core.py:158: UserWarning: [15:11:49]
WARNING: C:\buildkite-agent\builds\buildkite-windows-cpu-autoscaling-group-i-0c55ff5f71b100e98-1\xgboost\xgboost-ci-windows\src\learner.cc:740:
Parameters: { "use_label_encoder" } are not used.

warnings.warn(smsg, UserWarning)



Model saved as XGBClassifier_pipeline.joblib

Predictions on first 5 rows of test_variance using XGBClassifier:

[0 0 0 0 0]

🏃 View run XGBClassifier_Variance at: <https://dagshub.com/nithinyanna3/my-first-repo.mlflow/#/experiments/0/runs/a9122eb3c2a84bda89bdc0b3994da7cc>

🔧 View experiment at: <https://dagshub.com/nithinyanna3/my-first-repo.mlflow/#/experiments/0>

All models have been trained, logged, and saved successfully.

Use PCA for dimensionality reduction on all the features and testing it on random forest

In [244...

```
import pandas as pd
import mlflow
import mlflow.sklearn
from sklearn.decomposition import PCA
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, f1_score, confusion_matrix
from sklearn.preprocessing import StandardScaler
import numpy as np
import joblib

# Load dataset
df = secondupdate.copy()

# Separate features and target
X = df.drop("loan_status", axis=1)
y = df["loan_status"]

# Standardization
```

```
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, ran

# PCA
pca = PCA()
X_train_pca = pca.fit_transform(X_train)
X_test_pca = pca.transform(X_test)

# Determine number of components for 95% explained variance
explained_variance = pca.explained_variance_ratio_
cumulative_explained_variance = np.cumsum(explained_variance)
n_components = np.argmax(cumulative_explained_variance >= 0.95) + 1
print(f"Number of components selected for 95% explained variance: {n_components}")

# Apply PCA with the selected number of components
pca = PCA(n_components=n_components)
X_train_pca = pca.fit_transform(X_train)
X_test_pca = pca.transform(X_test)

# Train RandomForest model
model = RandomForestClassifier(random_state=42)
model.fit(X_train_pca, y_train)

# Save model using joblib
joblib.dump(model, "random_forest_pca_model.joblib")

mlflow.set_tracking_uri("https://dagshub.com/nithinyanna3/my-first-repo.mlflow") #
mlflow.set_experiment("PCA")

# Log model to MLflow
with mlflow.start_run():
    mlflow.sklearn.log_model(model, "RandomForest_PCA_Model")

# Predictions
y_pred = model.predict(X_test_pca)
acc = accuracy_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred, average='weighted')
cm = confusion_matrix(y_test, y_pred)

# Log metrics
mlflow.log_metric("accuracy", acc)
mlflow.log_metric("f1_score", f1)

# Confusion Matrix
tp, fn, fp, tn = cm.ravel()
mlflow.log_metric("TP", tp)
mlflow.log_metric("TN", tn)
mlflow.log_metric("FP", fp)
mlflow.log_metric("FN", fn)

# Print results
print("Random Forest with PCA Classification Metrics:")
print(f"Accuracy: {acc}")
```

```

print(f"F1 Score (weighted): {f1}")
print(f"Confusion Matrix:\n{cm}")
print(f"True Positive (TP): {tp}, False Negative (FN): {fn}, False Positive (FP)

# Test with 5 rows from updatedtestdata.csv
df_test = pd.read_csv("updatedtestdata.csv")
if "loan_status" in df_test.columns:
    X_test_sample = df_test.iloc[:5]
    y_test_sample = df_test["loan_status"].iloc[:5]
    X_test_sample_scaled = scaler.transform(X_test_sample)
    X_test_sample_pca = pca.transform(X_test_sample_scaled)
    y_test_sample_pred = model.predict(X_test_sample_pca)

    # Print 5 test samples predictions
    for i in range(5):
        print(f"Row {i + 1}: Predicted loan_status = {y_test_sample_pred[i]}, A
else:
    print("The 'loan_status' column is missing in the test dataset.")

print("Model training and logging completed.")

```

Number of components selected for 95% explained variance: 5

2024/12/17 15:39:17 WARNING mlflow.models.model: Model logged without a signature and input example. Please set `input_example` parameter when logging the model to auto infer the model signature.

Random Forest with PCA Classification Metrics:

Accuracy: 0.8995005167068549


F1 Score (weighted): 0.8908174486589133


Confusion Matrix:

```
[[9611 299]
 [ 868 834]]
```

True Positive (TP): 9611, False Negative (FN): 299, False Positive (FP): 868, True Negative (TN): 834

The 'loan_status' column is missing in the test dataset.

 View run charming-bass-298 at: <https://dagshub.com/nithinyanna3/my-first-repo.mlflow/#/experiments/5/runs/b656d71806fb4f22af40c29e6fb28b13>

 View experiment at: <https://dagshub.com/nithinyanna3/my-first-repo.mlflow/#/experiments/5>

Model training and logging completed.

the custom experiment i performed is KNN with PCA and logged into dagshub

In [248...

```

import pandas as pd
import mlflow
import mlflow.sklearn
from sklearn.decomposition import PCA
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, f1_score, confusion_matrix
from sklearn.preprocessing import StandardScaler
import numpy as np
import joblib

# Load dataset
df = secondupdate.copy()

```

```
# Separate features and target
X = df.drop("loan_status", axis=1)
y = df["loan_status"]

# Standardization
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, ran

# PCA
pca = PCA()
X_train_pca = pca.fit_transform(X_train)
X_test_pca = pca.transform(X_test)

# Determine number of components for 95% explained variance
explained_variance = pca.explained_variance_ratio_
cumulative_explained_variance = np.cumsum(explained_variance)
n_components = np.argmax(cumulative_explained_variance >= 0.95) + 1
print(f"Number of components selected for 95% explained variance: {n_components}")

# Apply PCA with the selected number of components
pca = PCA(n_components=n_components)
X_train_pca = pca.fit_transform(X_train)
X_test_pca = pca.transform(X_test)

# Train KNN model
model = KNeighborsClassifier(n_neighbors=5)
model.fit(X_train_pca, y_train)

# Save model using joblib
joblib.dump(model, "knn_pca_model.joblib")

mlflow.set_tracking_uri("https://dagshub.com/nithinyanna3/my-first-repo.mlflow") #
mlflow.set_experiment("PCA")

# Log model to MLflow
with mlflow.start_run():
    mlflow.sklearn.log_model(model, "KNN_PCA_Model")

# Predictions
y_pred = model.predict(X_test_pca)
acc = accuracy_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred, average='weighted')
cm = confusion_matrix(y_test, y_pred)

# Log metrics
mlflow.log_metric("accuracy", acc)
mlflow.log_metric("f1_score", f1)

# Confusion Matrix
tp, fn, fp, tn = cm.ravel()
mlflow.log_metric("TP", tp)
mlflow.log_metric("TN", tn)
mlflow.log_metric("FP", fp)
```

```

mlflow.log_metric("FN", fn)

# Print results
print("KNN with PCA Classification Metrics:")
print(f"Accuracy: {acc}")
print(f"F1 Score (weighted): {f1}")
print(f"Confusion Matrix:\n{cm}")
print(f"True Positive (TP): {tp}, False Negative (FN): {fn}, False Positive (FP)
# Test with 5 rows from updatedtestdata.csv
df_test = pd.read_csv("updatedtestdata.csv")
if "loan_status" in df_test.columns:
    X_test_sample = df_test.iloc[:5]
    y_test_sample = df_test["loan_status"].iloc[:5]
    X_test_sample_scaled = scaler.transform(X_test_sample)
    X_test_sample_pca = pca.transform(X_test_sample_scaled)
    y_test_sample_pred = model.predict(X_test_sample_pca)

    # Print 5 test samples predictions
    for i in range(5):
        print(f"Row {i + 1}: Predicted loan_status = {y_test_sample_pred[i]}, A
else:
    print("The 'loan_status' column is missing in the test dataset.")

print("Model training and logging completed.")

```

Number of components selected for 95% explained variance: 5

2024/12/17 15:47:13 WARNING mlflow.models.model: Model logged without a signature and input example. Please set `input_example` parameter when logging the model to auto infer the model signature.

KNN with PCA Classification Metrics:

Accuracy: 0.8944195659662418


F1 Score (weighted): 0.8865163979822601


Confusion Matrix:

```
[[9550  360]
 [ 866  836]]
```

True Positive (TP): 9550, False Negative (FN): 360, False Positive (FP): 866, True Negative (TN): 836

The 'loan_status' column is missing in the test dataset.

 View run nosy-fowl-4 at: <https://dagshub.com/nithinyanna3/my-first-repo.mlflow/#/experiments/5/runs/2ebe3c8c0a914053b6c42cf1b4ccd2ef>

 View experiment at: <https://dagshub.com/nithinyanna3/my-first-repo.mlflow/#/experiments/5>

Model training and logging completed.

the second custom experiment i performed is SVM with PCA

In [253...

```

import pandas as pd
import mlflow
import mlflow.sklearn
from sklearn.decomposition import PCA
from sklearn.model_selection import train_test_split
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, f1_score, confusion_matrix
from sklearn.preprocessing import StandardScaler
import numpy as np
import joblib

```



```
# Load dataset
df = secondupdate.copy()

# Separate features and target
X = df.drop("loan_status", axis=1)
y = df["loan_status"]

# Standardization
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, ran

# PCA
pca = PCA()
X_train_pca = pca.fit_transform(X_train)
X_test_pca = pca.transform(X_test)

# Determine number of components for 95% explained variance
explained_variance = pca.explained_variance_ratio_
cumulative_explained_variance = np.cumsum(explained_variance)
n_components = np.argmax(cumulative_explained_variance >= 0.95) + 1
print(f"Number of components selected for 95% explained variance: {n_components}")

# Apply PCA with the selected number of components
pca = PCA(n_components=n_components)
X_train_pca = pca.fit_transform(X_train)
X_test_pca = pca.transform(X_test)

# Train SVM model
model = SVC(kernel='rbf', random_state=42)
model.fit(X_train_pca, y_train)

# Save model using joblib
joblib.dump(model, "svm_pca_model.joblib")

mlflow.set_tracking_uri("https://dagshub.com/nithinyanna3/my-first-repo.mlflow") #
mlflow.set_experiment("PCA")

# Log model to MLflow
with mlflow.start_run():
    mlflow.sklearn.log_model(model, "SVM_PCA_Model")

# Predictions
y_pred = model.predict(X_test_pca)
acc = accuracy_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred, average='weighted')
cm = confusion_matrix(y_test, y_pred)

# Log metrics
mlflow.log_metric("accuracy", acc)
mlflow.log_metric("f1_score", f1)

# Confusion Matrix
```

```

tp, fn, fp, tn = cm.ravel()
mlflow.log_metric("TP", tp)
mlflow.log_metric("TN", tn)
mlflow.log_metric("FP", fp)
mlflow.log_metric("FN", fn)

# Print results
print("SVM with PCA Classification Metrics:")
print(f"Accuracy: {acc}")
print(f"F1 Score (weighted): {f1}")
print(f"Confusion Matrix:\n{cm}")
print(f"True Positive (TP): {tp}, False Negative (FN): {fn}, False Positive (FP)

# Test with 5 rows from updatedtestdata.csv
df_test = pd.read_csv("updatedtestdata.csv")
if "loan_status" in df_test.columns:
    X_test_sample = df_test.iloc[:5]
    y_test_sample = df_test["loan_status"].iloc[:5]
    X_test_sample_scaled = scaler.transform(X_test_sample)
    X_test_sample_pca = pca.transform(X_test_sample_scaled)
    y_test_sample_pred = model.predict(X_test_sample_pca)

    # Print 5 test samples predictions
    for i in range(5):
        print(f"Row {i + 1}: Predicted loan_status = {y_test_sample_pred[i]}, A
else:
    print("The 'loan_status' column is missing in the test dataset.")

print("Model training and logging completed.")

```

Number of components selected for 95% explained variance: 5

2024/12/17 15:55:43 WARNING mlflow.models.model: Model logged without a signature and input example. Please set `input_example` parameter when logging the model to auto infer the model signature.

SVM with PCA Classification Metrics:

Accuracy: 0.8965725111953152

F1 Score (weighted): 0.8827149247633431


Confusion Matrix:


[[9709 201]

[1000 702]]

True Positive (TP): 9709, False Negative (FN): 201, False Positive (FP): 1000, True Negative (TN): 702

The 'loan_status' column is missing in the test dataset.

 View run PCA with SVM at: <https://dagshub.com/nithinyanna3/my-first-repo.mlflow/#/experiments/5/runs/5790d03e0dc04550885c47f049ad18f1>

 View experiment at: <https://dagshub.com/nithinyanna3/my-first-repo.mlflow/#/experiments/5>

Model training and logging completed.

Create meaningful F1-score plots and accuracies to compare experiments and determine the best model.

In []:

printing output from best model with 92% accuracy XGBoost

```
In [267... import pandas as pd
import mlflow
import xgboost as xgb
import joblib

# Load the saved XGBoost model
model = joblib.load("xgb_classifier_model.pkl")

# Load the test dataset
df_test = pd.read_csv("updatedtestdata.csv")

# Predict
y_test_sample_pred = model.predict(df_test)

# Add predictions to the DataFrame
df_test['Predicted loan_status'] = y_test_sample_pred

# Display the test data with predictions
print(df_test)

print("Prediction on test data completed.")
```

	person_age	person_income	person_emp_length	loan_grade	loan_amnt	\
0	23	69000	3.0	5.0	25000	
1	26	96000	6.0	2.0	10000	
2	26	30000	5.0	4.0	4000	
3	33	50000	4.0	0.0	7000	
4	26	102000	8.0	3.0	15000	
...	
39093	22	31200	2.0	1.0	3000	
39094	22	48000	6.0	0.0	7000	
39095	51	60000	0.0	0.0	15000	
39096	22	36000	4.0	3.0	14000	
39097	31	45000	6.0	1.0	19450	

	loan_int_rate	loan_percent_income	cb_person_default_on_file	\
0	15.76	0.36	1	
1	12.68	0.10	2	
2	17.19	0.13	2	
3	8.90	0.14	1	
4	16.32	0.15	2	
...	
39093	10.37	0.10	1	
39094	6.03	0.15	1	
39095	7.51	0.25	1	
39096	15.62	0.39	2	
39097	9.91	0.44	1	

	cb_person_cred_hist_length	Predicted loan_status
0	2	1
1	4	0
2	2	1
3	7	0
4	4	0
...
39093	4	0
39094	3	0
39095	25	0
39096	4	1
39097	9	1

[39098 rows x 10 columns]

Prediction on test data completed.

Save the final model using joblib. Create a FastAPI application to serve the model.

Containerize the FastAPI application using Docker and push to Docker Hub. Deploy the containerized API to a cloud platform.

Create a Streamlit app to interact with the deployed model for real-time classification.

In [263... `pip install scikit-learn joblib`

Requirement already satisfied: scikit-learn in c:\users\hp\anaconda3\lib\site-packages (1.4.2)
 Requirement already satisfied: joblib in c:\users\hp\anaconda3\lib\site-packages (1.4.2)
 Requirement already satisfied: numpy>=1.19.5 in c:\users\hp\anaconda3\lib\site-packages (from scikit-learn) (1.26.4)
 Requirement already satisfied: scipy>=1.6.0 in c:\users\hp\anaconda3\lib\site-packages (from scikit-learn) (1.13.1)
 Requirement already satisfied: threadpoolctl>=2.0.0 in c:\users\hp\anaconda3\lib\site-packages (from scikit-learn) (2.2.0)
 Note: you may need to restart the kernel to use updated packages.

In [265...

```
import joblib
from xgboost import XGBClassifier

# Assuming 'model' is your trained XGBoost model
joblib.dump(model, 'xgb_classifier_model.pkl')
```

Out[265...

```
['xgb_classifier_model.pkl']
```

In [273...

```
pip install fastapi uvicorn
```

Requirement already satisfied: fastapi in c:\users\hp\anaconda3\lib\site-packages (0.115.6)
 Requirement already satisfied: uvicorn in c:\users\hp\anaconda3\lib\site-packages (0.34.0)
 Requirement already satisfied: starlette<0.42.0,>=0.40.0 in c:\users\hp\anaconda3\lib\site-packages (from fastapi) (0.41.3)
 Requirement already satisfied: pydantic!=1.8,!<1.8.1,!<2.0.0,!<2.0.1,!<2.1.0,<3.0.0,>=1.7.4 in c:\users\hp\anaconda3\lib\site-packages (from fastapi) (2.5.3)
 Requirement already satisfied: typing-extensions>=4.8.0 in c:\users\hp\anaconda3\lib\site-packages (from fastapi) (4.11.0)
 Requirement already satisfied: click>=7.0 in c:\users\hp\anaconda3\lib\site-packages (from uvicorn) (8.1.7)
 Requirement already satisfied: h11>=0.8 in c:\users\hp\anaconda3\lib\site-packages (from uvicorn) (0.14.0)
 Requirement already satisfied: colorama in c:\users\hp\anaconda3\lib\site-packages (from click>=7.0->uvicorn) (0.4.6)
 Requirement already satisfied: annotated-types>=0.4.0 in c:\users\hp\anaconda3\lib\site-packages (from pydantic!=1.8,!<1.8.1,!<2.0.0,!<2.0.1,!<2.1.0,<3.0.0,>=1.7.4->fastapi) (0.6.0)
 Requirement already satisfied: pydantic-core==2.14.6 in c:\users\hp\anaconda3\lib\site-packages (from pydantic!=1.8,!<1.8.1,!<2.0.0,!<2.0.1,!<2.1.0,<3.0.0,>=1.7.4->fastapi) (2.14.6)
 Requirement already satisfied: anyio<5,>=3.4.0 in c:\users\hp\anaconda3\lib\site-packages (from starlette<0.42.0,>=0.40.0->fastapi) (4.2.0)
 Requirement already satisfied: idna>=2.8 in c:\users\hp\anaconda3\lib\site-packages (from anyio<5,>=3.4.0->starlette<0.42.0,>=0.40.0->fastapi) (3.7)
 Requirement already satisfied: sniffio>=1.1 in c:\users\hp\anaconda3\lib\site-packages (from anyio<5,>=3.4.0->starlette<0.42.0,>=0.40.0->fastapi) (1.3.0)
 Note: you may need to restart the kernel to use updated packages.

local deployment of fastapi

In [4]:

```
import uvicorn
from fastapi import FastAPI
```

```

from pydantic import BaseModel
import numpy as np
import pandas as pd
import joblib
import nest_asyncio

# Apply nest_asyncio to fix event loop issues in Jupyter
nest_asyncio.apply()

# Load the model
model = joblib.load("best_xgb_model_with_scaler.pkl")

app = FastAPI()

class InputData(BaseModel):
    person_age: float
    person_income: float
    person_emp_length: float
    loan_grade: float
    loan_amnt: float
    loan_int_rate: float
    loan_percent_income: float
    cb_person_default_on_file: float
    cb_person_cred_hist_length: float

@app.post("/predict")
async def predict(data: InputData):
    try:
        input_features = np.array([
            data.person_age,
            data.person_income,
            data.person_emp_length,
            data.loan_grade,
            data.loan_amnt,
            data.loan_int_rate,
            data.loan_percent_income,
            data.cb_person_default_on_file,
            data.cb_person_cred_hist_length,
        ])

        # Convert the array to a DataFrame with the appropriate column names
        input_df = pd.DataFrame([input_features], columns=[
            'person_age', 'person_income', 'person_emp_length', 'loan_grade',
            'loan_amnt', 'loan_int_rate', 'loan_percent_income',
            'cb_person_default_on_file', 'cb_person_cred_hist_length'
        ])

        # Make prediction
        prediction = model.predict(input_df)
        return {"prediction": int(prediction[0])}
    except Exception as e:
        return {"error": str(e)}

if __name__ == "__main__":
    config = uvicorn.Config(app, host="127.0.0.1", port=8002, log_level="info")

```

```
server = uvicorn.Server(config)
await server.serve()
```

```
INFO:      Started server process [14084]
INFO:      Waiting for application startup.
INFO:      Application startup complete.
INFO:      Uvicorn running on http://127.0.0.1:8002 (Press CTRL+C to quit)
INFO:      Shutting down
INFO:      Waiting for application shutdown.
INFO:      Application shutdown complete.
INFO:      Finished server process [14084]
```

In [6]: `!pip install docker render-python`

```
Requirement already satisfied: docker in c:\users\hp\anaconda3\lib\site-packages (7.1.0)
Collecting render-python
  Downloading render_python-2.3.1-py3-none-any.whl.metadata (567 bytes)
Requirement already satisfied: pywin32>=304 in c:\users\hp\anaconda3\lib\site-packages (from docker) (305.1)
Requirement already satisfied: requests>=2.26.0 in c:\users\hp\anaconda3\lib\site-packages (from docker) (2.32.2)
Requirement already satisfied: urllib3>=1.26.0 in c:\users\hp\anaconda3\lib\site-packages (from docker) (2.2.2)
Requirement already satisfied: decorator in c:\users\hp\anaconda3\lib\site-packages (from render-python) (5.1.1)
Requirement already satisfied: numpy in c:\users\hp\anaconda3\lib\site-packages (from render-python) (1.26.4)
Requirement already satisfied: pillow in c:\users\hp\anaconda3\lib\site-packages (from render-python) (10.3.0)
Requirement already satisfied: scipy in c:\users\hp\anaconda3\lib\site-packages (from render-python) (1.13.1)
Requirement already satisfied: six in c:\users\hp\anaconda3\lib\site-packages (from render-python) (1.16.0)
Collecting sphinxcontrib-napoleon (from render-python)
  Downloading sphinxcontrib_napoleon-0.7-py2.py3-none-any.whl.metadata (6.2 kB)
Requirement already satisfied: charset-normalizer<4,>=2 in c:\users\hp\anaconda3\lib\site-packages (from requests>=2.26.0->docker) (2.0.4)
Requirement already satisfied: idna<4,>=2.5 in c:\users\hp\anaconda3\lib\site-packages (from requests>=2.26.0->docker) (3.7)
Requirement already satisfied: certifi>=2017.4.17 in c:\users\hp\anaconda3\lib\site-packages (from requests>=2.26.0->docker) (2024.8.30)
Collecting pockets>=0.3 (from sphinxcontrib-napoleon->render-python)
  Downloading pockets-0.9.1-py2.py3-none-any.whl.metadata (4.4 kB)
Downloading render_python-2.3.1-py3-none-any.whl (67 kB)
Downloading sphinxcontrib_napoleon-0.7-py2.py3-none-any.whl (17 kB)
Downloading pockets-0.9.1-py2.py3-none-any.whl (26 kB)
Installing collected packages: pockets, sphinxcontrib-napoleon, render-python
Successfully installed pockets-0.9.1 render-python-2.3.1 sphinxcontrib-napoleon-0.7
```

working with dockerfile directory

In [3]: `dockerfile_content = """
Use the official Python image as the base image
FROM python:3.10-slim`

```
# Set the working directory in the container
WORKDIR /app

# Copy the requirements file to the container
COPY requirements.txt .

# Install the required Python libraries
RUN pip install --no-cache-dir -r requirements.txt

# Copy the FastAPI app code to the container
COPY . .

# Expose the application port
EXPOSE 8001

# Command to run the application
CMD ["uvicorn", "main:app", "--host", "0.0.0.0", "--port", "8002"]
"""

with open("Dockerfile", "w") as file:
    file.write(dockerfile_content)
```

```
In [5]: requirements_content = """
fastapi
uvicorn
numpy
pandas
joblib
nest_asyncio
"""

with open("requirements.txt", "w") as file:
    file.write(requirements_content)
```

```
In [11]: import os

os.chdir("C:\\Users\\HP")
```

```
In [17]: mkdir C:\\Users\\HP\\docker_build
```

```
In [13]: import uvicorn
from fastapi import FastAPI
from pydantic import BaseModel
import numpy as np
import pandas as pd
import joblib
import os

# Define the Python code as a string
python_code = """
import uvicorn
from fastapi import FastAPI
from pydantic import BaseModel
import numpy as np
import pandas as pd
```



```
import joblib

# Load the model
model = joblib.load("best_xgb_model_with_scaler.pkl")

# Initialize the FastAPI app
app = FastAPI()

# Define the input data schema
class InputData(BaseModel):
    person_age: float
    person_income: float
    person_emp_length: float
    loan_grade: float
    loan_amnt: float
    loan_int_rate: float
    loan_percent_income: float
    cb_person_default_on_file: float
    cb_person_cred_hist_length: float

# Prediction endpoint
@app.post("/predict")
async def predict(data: InputData):
    try:
        # Prepare the input features for the model
        input_features = np.array([
            data.person_age,
            data.person_income,
            data.person_emp_length,
            data.loan_grade,
            data.loan_amnt,
            data.loan_int_rate,
            data.loan_percent_income,
            data.cb_person_default_on_file,
            data.cb_person_cred_hist_length,
        ])

        # Convert to a DataFrame
        input_df = pd.DataFrame([input_features], columns=[
            'person_age', 'person_income', 'person_emp_length', 'loan_grade',
            'loan_amnt', 'loan_int_rate', 'loan_percent_income',
            'cb_person_default_on_file', 'cb_person_cred_hist_length'
        ])

        # Make a prediction
        prediction = model.predict(input_df)
        return {"prediction": int(prediction[0])}
    except Exception as e:
        # Return error details if something goes wrong
        return {"error": str(e)}

if __name__ == "__main__":
    uvicorn.run(app, host="0.0.0.0", port=8002, log_level="info")
"""

# Define the target directory
```

```
target_dir = r"C:\Users\HP\docker_build"

# Create the target directory if it doesn't exist
os.makedirs(target_dir, exist_ok=True)

# Save the script as main.py in the target directory
file_path = os.path.join(target_dir, "main.py")
try:
    with open(file_path, "w") as file:
        file.write(python_code)
    print(f"main.py saved successfully in {target_dir}")
except Exception as e:
    print(f"Error saving main.py: {e}")
```

main.py saved successfully in C:\Users\HP\docker_build

```
In [21]: import os
import shutil

# Create target directory if it doesn't exist
target_dir = r"C:\Users\HP\docker_build"
os.makedirs(target_dir, exist_ok=True)

# Move the Dockerfile
shutil.move(r"C:\Users\HP\Dockerfile", target_dir)

# Move other necessary files (example: requirements.txt)
shutil.move(r"C:\Users\HP\requirements.txt", target_dir)
```

Out[21]: 'C:\\Users\\HP\\docker_build\\requirements.txt'

```
In [17]: import docker

# Initialize Docker client
client = docker.from_env()

# Build the Docker image
try:
    image, build_logs = client.images.build(
        path=r"C:\Users\HP\docker_build",
        tag="nithinyanna/fastapi-app:v1"
    )
    print("Docker image built successfully!")
except Exception as e:
    print(f"Error building Docker image: {e}")
```

```
Task exception was never retrieved
future: <Task finished name='Task-1' coro=<Server.serve() done, defined at C:\Users\
\HP\anaconda3\Lib\site-packages\uvicorn\server.py:68> exception=KeyboardInterrupt()>
Traceback (most recent call last):
  File "C:\Users\HP\anaconda3\Lib\site-packages\uvicorn\main.py", line 579, in run
    server.run()
  File "C:\Users\HP\anaconda3\Lib\site-packages\uvicorn\server.py", line 66, in run
    return asyncio.run(self.serve(sockets=sockets))
    ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
  File "C:\Users\HP\anaconda3\Lib\site-packages\nest_asyncio.py", line 30, in run
    return loop.run_until_complete(task)
    ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
  File "C:\Users\HP\anaconda3\Lib\site-packages\nest_asyncio.py", line 92, in run_un
til_complete
    self._run_once()
  File "C:\Users\HP\anaconda3\Lib\site-packages\nest_asyncio.py", line 133, in _run_
once
    handle._run()
  File "C:\Users\HP\anaconda3\Lib\asyncio\events.py", line 88, in _run
    self._context.run(self._callback, *self._args)
  File "C:\Users\HP\anaconda3\Lib\asyncio\tasks.py", line 396, in __wakeup
    self.__step()
  File "C:\Users\HP\anaconda3\Lib\asyncio\tasks.py", line 303, in __step
    self.__step_run_and_handle_result(exc)
  File "C:\Users\HP\anaconda3\Lib\asyncio\tasks.py", line 314, in __step_run_and_han
dle_result
    result = coro.send(None)
    ^^^^^^^^^^^^^^^^^^^^^^^^^
  File "C:\Users\HP\anaconda3\Lib\site-packages\uvicorn\server.py", line 69, in serv
e
    with self.capture_signals():
  File "C:\Users\HP\anaconda3\Lib\contextlib.py", line 144, in __exit__
    next(self.gen)
  File "C:\Users\HP\anaconda3\Lib\site-packages\uvicorn\server.py", line 330, in cap
ture_signals
    signal.raise_signal(captured_signal)
KeyboardInterrupt
C:\Users\HP\anaconda3\Lib\site-packages\paramiko\transport.py:219: CryptographyDepre
cationWarning: Blowfish has been deprecated and will be removed in a future release
  "class": algorithms.Blowfish,
Docker image built successfully!
```

```
In [25]: import docker

client = docker.from_env()

try:
    image, logs = client.images.build(path=r"C:\Users\HP\docker_build", tag="nithin")
    print("Docker image built successfully!")
    print(f"Image ID: {image.id}")
    # Print build logs (optional)
    for log in logs:
        if 'stream' in log:
            print(log['stream'].strip())
except docker.errors.BuildError as e:
    print("Error during build:", e)
```

```
except docker.errors.APIError as e:
    print("Docker API error:", e)
```

Docker image built successfully!

Image ID: sha256:24ce2b9269b6b9760d33bf0e03423fdb99f567b49dd8e094b79535285042fd54

Step 1/7 : FROM python:3.10-slim

---> 61912260e578

Step 2/7 : WORKDIR /app

---> Using cache

---> 4496453a7ea0

Step 3/7 : COPY requirements.txt .

---> Using cache

---> aa7de811e07f

Step 4/7 : RUN pip install --no-cache-dir -r requirements.txt

---> Using cache

---> e7a5aee47fc

Step 5/7 : COPY . .

---> Using cache

---> 1f55e8daf29f

Step 6/7 : EXPOSE 8001

---> Using cache

---> e59cb9fd0778

Step 7/7 : CMD ["uvicorn", "main:app", "--host", "0.0.0.0", "--port", "8002"]

---> Using cache

---> 24ce2b9269b6

Successfully built 24ce2b9269b6

Successfully tagged nithinyanna/fastapi-app:v1

successfully created dockerimages

```
In [19]: client.login(username="nithinyanna", password="N@ithin789")
client.images.push("nithinyanna/fastapi-app:v1")
print("Docker image pushed to Docker Hub.")
```

Docker image pushed to Docker Hub.

to deploy the model using cloud i have used render cloud which is easy and reliable and below is the loal deplyment of streamlitt application

```
In [1]: pip install streamlit
```

Requirement already satisfied: streamlit in c:\users\hp\anaconda3\lib\site-packages (1.32.0)

Requirement already satisfied: altair<6,>=4.0 in c:\users\hp\anaconda3\lib\site-packages (from streamlit) (5.0.1)

Requirement already satisfied: blinker<2,>=1.0.0 in c:\users\hp\anaconda3\lib\site-packages (from streamlit) (1.6.2)

Requirement already satisfied: cachetools<6,>=4.0 in c:\users\hp\anaconda3\lib\site-packages (from streamlit) (5.3.3)

Requirement already satisfied: click<9,>=7.0 in c:\users\hp\anaconda3\lib\site-packages (from streamlit) (8.1.7)

Requirement already satisfied: numpy<2,>=1.19.3 in c:\users\hp\anaconda3\lib\site-packages (from streamlit) (1.26.4)

Requirement already satisfied: packaging<24,>=16.8 in c:\users\hp\anaconda3\lib\site-packages (from streamlit) (23.2)

Requirement already satisfied: pandas<3,>=1.3.0 in c:\users\hp\anaconda3\lib\site-packages (from streamlit) (2.2.2)

Requirement already satisfied: pillow<11,>=7.1.0 in c:\users\hp\anaconda3\lib\site-packages (from streamlit) (10.3.0)

Requirement already satisfied: protobuf<5,>=3.20 in c:\users\hp\anaconda3\lib\site-packages (from streamlit) (3.20.3)

Requirement already satisfied: pyarrow>=7.0 in c:\users\hp\anaconda3\lib\site-packages (from streamlit) (14.0.2)

Requirement already satisfied: requests<3,>=2.27 in c:\users\hp\anaconda3\lib\site-packages (from streamlit) (2.32.2)

Requirement already satisfied: rich<14,>=10.14.0 in c:\users\hp\anaconda3\lib\site-packages (from streamlit) (13.3.5)

Requirement already satisfied: tenacity<9,>=8.1.0 in c:\users\hp\anaconda3\lib\site-packages (from streamlit) (8.2.2)

Requirement already satisfied: toml<2,>=0.10.1 in c:\users\hp\anaconda3\lib\site-packages (from streamlit) (0.10.2)

Requirement already satisfied: typing-extensions<5,>=4.3.0 in c:\users\hp\anaconda3\lib\site-packages (from streamlit) (4.11.0)

Requirement already satisfied: gitpython!=3.1.19,<4,>=3.0.7 in c:\users\hp\anaconda3\lib\site-packages (from streamlit) (3.1.37)

Requirement already satisfied: pydeck<1,>=0.8.0b4 in c:\users\hp\anaconda3\lib\site-packages (from streamlit) (0.8.0)

Requirement already satisfied: tornado<7,>=6.0.3 in c:\users\hp\anaconda3\lib\site-packages (from streamlit) (6.4.1)

Requirement already satisfied: watchdog>=2.1.5 in c:\users\hp\anaconda3\lib\site-packages (from streamlit) (4.0.1)

Requirement already satisfied: jinja2 in c:\users\hp\anaconda3\lib\site-packages (from altair<6,>=4.0->streamlit) (3.1.4)

Requirement already satisfied: jsonschema>=3.0 in c:\users\hp\anaconda3\lib\site-packages (from altair<6,>=4.0->streamlit) (4.19.2)

Requirement already satisfied: toolz in c:\users\hp\anaconda3\lib\site-packages (from altair<6,>=4.0->streamlit) (0.12.0)

Requirement already satisfied: colorama in c:\users\hp\anaconda3\lib\site-packages (from click<9,>=7.0->streamlit) (0.4.6)

Requirement already satisfied: gitdb<5,>=4.0.1 in c:\users\hp\anaconda3\lib\site-packages (from gitpython!=3.1.19,<4,>=3.0.7->streamlit) (4.0.7)

Requirement already satisfied: python-dateutil>=2.8.2 in c:\users\hp\anaconda3\lib\site-packages (from pandas<3,>=1.3.0->streamlit) (2.9.0.post0)

Requirement already satisfied: pytz>=2020.1 in c:\users\hp\anaconda3\lib\site-packages (from pandas<3,>=1.3.0->streamlit) (2024.1)

Requirement already satisfied: tzdata>=2022.7 in c:\users\hp\anaconda3\lib\site-packages (from pandas<3,>=1.3.0->streamlit) (2023.3)

Requirement already satisfied: charset-normalizer<4,>=2 in c:\users\hp\anaconda3\lib\site-packages (from requests<3,>=2.27->streamlit) (2.0.4)

Requirement already satisfied: idna<4,>=2.5 in c:\users\hp\anaconda3\lib\site-packages (from requests<3,>=2.27->streamlit) (3.7)

Requirement already satisfied: urllib3<3,>=1.21.1 in c:\users\hp\anaconda3\lib\site-packages (from requests<3,>=2.27->streamlit) (2.2.2)

Requirement already satisfied: certifi>=2017.4.17 in c:\users\hp\anaconda3\lib\site-packages (from requests<3,>=2.27->streamlit) (2024.8.30)

Requirement already satisfied: markdown-it-py<3.0.0,>=2.2.0 in c:\users\hp\anaconda3\lib\site-packages (from rich<14,>=10.14.0->streamlit) (2.2.0)

Requirement already satisfied: pygments<3.0.0,>=2.13.0 in c:\users\hp\anaconda3\lib\site-packages (from rich<14,>=10.14.0->streamlit) (2.15.1)

Requirement already satisfied: smmap<5,>=3.0.1 in c:\users\hp\anaconda3\lib\site-packages (from gitdb<5,>=4.0.1->gitpython!=3.1.19,<4,>=3.0.7->streamlit) (4.0.0)

Requirement already satisfied: MarkupSafe>=2.0 in c:\users\hp\anaconda3\lib\site-packages (from jinja2->altair<6,>=4.0->streamlit) (2.1.3)

Requirement already satisfied: attrs>=22.2.0 in c:\users\hp\anaconda3\lib\site-packages (from jsonschema>=3.0->altair<6,>=4.0->streamlit) (23.1.0)

Requirement already satisfied: jsonschema-specifications>=2023.03.6 in c:\users\hp\anaconda3\lib\site-packages (from jsonschema>=3.0->altair<6,>=4.0->streamlit) (2023.7.1)

Requirement already satisfied: referencing>=0.28.4 in c:\users\hp\anaconda3\lib\site-packages (from jsonschema>=3.0->altair<6,>=4.0->streamlit) (0.30.2)

Requirement already satisfied: rpds-py>=0.7.1 in c:\users\hp\anaconda3\lib\site-packages (from jsonschema>=3.0->altair<6,>=4.0->streamlit) (0.10.6)

Requirement already satisfied: mdurl~0.1 in c:\users\hp\anaconda3\lib\site-packages (from markdown-it-py<3.0.0,>=2.2.0->rich<14,>=10.14.0->streamlit) (0.1.0)

Requirement already satisfied: six>=1.5 in c:\users\hp\anaconda3\lib\site-packages (from python-dateutil>=2.8.2->pandas<3,>=1.3.0->streamlit) (1.16.0)

Note: you may need to restart the kernel to use updated packages.

```
In [ ]: import streamlit as st
import joblib
import pandas as pd
import numpy as np

# Load the model
model = joblib.load("best_xgb_model_with_scaler.pkl")

# Streamlit app
st.title("Loan Default Prediction")

# Define input fields
person_age = st.number_input("Age", min_value=18, max_value=100, value=30)
person_income = st.number_input("Income", min_value=0, value=50000)
person_emp_length = st.number_input("Employment Length (in years)", min_value=0, max_value=10, value=5)
loan_grade = st.number_input("Loan Grade", min_value=1, max_value=10, value=5)
loan_amnt = st.number_input("Loan Amount", min_value=0, value=10000)
loan_int_rate = st.number_input("Interest Rate (%)", min_value=0.0, value=5.0)
loan_percent_income = st.number_input("Loan % of Income", min_value=0.0, value=10.0)
cb_person_default_on_file = st.number_input("Default on File", min_value=0, max_value=1, value=0)
cb_person_cred_hist_length = st.number_input("Credit History Length", min_value=0, max_value=10, value=1)

# Create a DataFrame with the input data
input_data = pd.DataFrame([person_age, person_income, person_emp_length, loan_grade, loan_amnt, loan_int_rate, loan_percent_income, cb_person_default_on_file, cb_person_cred_hist_length])
```

```
        cb_person_default_on_file, cb_person_cred_hist_length]]
columns=['person_age', 'person_income', 'person_emp_length',
         'loan_amnt', 'loan_int_rate', 'loan_percent_income',
         'cb_person_default_on_file', 'cb_person_cred_hist_length']

# Predict button
if st.button('Predict'):
    prediction = model.predict(input_data)
    st.write(f"Prediction: {'Default' if prediction[0] == 1 else 'No Default'}")
```